Exploiting Causal Graph Priors with Posterior Sampling for Reinforcement Learning

Mirco Mutti[*]
Politecnico di Milano

Riccardo De Santi
ETH Zurich

Marcello Restelli
Politecnico di Milano

Alexander Marx[†]
ETH AI Center

Giorgia Ramponi[†]
ETH AI Center

Abstract

Posterior sampling allows the exploitation of prior knowledge of the environment's transition dynamics to improve the sample efficiency of reinforcement learning. The prior is typically specified as a class of parametric distributions, a task that can be cumbersome in practice, often resulting in the choice of uninformative priors. In this work, we propose a novel posterior sampling approach in which the prior is given as a (partial) causal graph over the environment's variables. The latter is often more natural to design, such as listing known causal dependencies between biometric features in a medical treatment study. Specifically, we propose a hierarchical Bayesian procedure, called C-PSRL, simultaneously learning the full causal graph at the higher level and the parameters of the resulting factored dynamics at the lower level. For this procedure, we provide an analysis of its Bayesian regret, which explicitly connects the regret rate with the degree of prior knowledge. Our numerical evaluation conducted in illustrative domains confirms that C-PSRL strongly improves the efficiency of posterior sampling with an uninformative prior while performing close to posterior sampling with the full causal graph.

1 Introduction

Posterior sampling [Thompson, 1933], a.k.a. Thompson sampling, is a powerful alternative to classic optimistic methods for Reinforcement Learning [RL, Sutton and Barto, 2018], as it guarantees outstanding sample efficiency [Osband et al., 2013] through an explicit model of the epistemic uncertainty that allows exploiting prior knowledge over the environment's dynamics. Specifically, Posterior Sampling for Reinforcement Learning [PSRL, Strens, 2000, Osband et al., 2013] implements a Bayesian procedure in which, at every episode k, (1) a model of the environment's dynamics is sampled from a parametric prior distribution P_k, (2) an optimal policy π_k is computed (e.g., through value iteration [Bellman, 1957]) according to the sampled model, (3) a posterior update is performed on the prior parameters to incorporate in P_{k+1} the evidence collected by running π_k in the true environment. Under the assumption that the true environment's dynamics are sampled with positive probability from the prior P_0, the latter procedure is provably efficient as it showcases a Bayesian regret that scales with $O(\sqrt{K})$ being K the total number of episodes [Osband and Van Roy, 2017].

Although posterior sampling has been also praised for its empirical prowess [Chapelle and Li, 2011], specifying the prior through a class of parametric distributions, a crucial requirement of PSRL, can be cumbersome in practice. Let us take a Dynamic Treatment Regime [DTR, Murphy, 2003] as an illustrative application. Here, we aim to overcome a patient's disease by choosing, at each stage, a treatment based on the patient's evolving conditions and previously

---

[*]Correspondence to <mirco.mutti@polimi.it>. [†]These authors co-supervised the project.

administered treatments. The goal is to identify the best treatment for the specific patient quickly. Medicine provides plenty of prior knowledge to help solve the DTR problem. However, it is not easy to translate this knowledge into a parametric prior distribution that is general enough to include the model of any patient while being sufficiently narrow to foster efficiency. Instead, it is remarkably easy to list some known causal relationships between patient's state variables, such as heart rate and blood pressure, or diabetes and glucose level. Those causal edges might come from experts' knowledge (e.g., physicians) or previous clinical studies. A prior in the form of a causal graph is more natural to specify for practitioners, who might be unaware of the intricacies of Bayesian statistics. Posterior sampling does not currently support the specification of the prior through a causal graph, which limits its applicability.

This paper proposes a novel posterior sampling methodology that can exploit a prior specified through a partial causal graph over the environment's variables. Notably, a complete causal graph allows for a factorization of the environment's dynamics, which can be then expressed as a Factored Markov Decision Process [FMDP, Boutilier et al., 2000]. PSRL can be applied to FMDPs, as demonstrated by previous work [Osband and Van Roy, 2014], where the authors assume to know the complete causal graph. However, this assumption is often unreasonable in practical applications.[1]

Instead, we assume to have partial knowledge of the causal graph, which leads to considering a set of plausible FMDPs. Taking inspiration from [Hong et al., 2020, 2022b,a, Kveton et al., 2021], we design a hierarchical Bayesian procedure, called **C**ausal **PSRL** (C-PSRL), extending PSRL to the setting where the true model lies within a set of FMDPs (induced by the causal graph prior). At each episode, C-PSRL first samples a factorization consistent with the causal graph prior. Then, it samples the model of the FMDP from a lower-level prior that is conditioned on the sampled factorization. After that, the algorithm proceeds similarly to PSRL on the sampled FMDP.

Having introduced C-PSRL, we study the Bayesian regret it induces on the footsteps of previous analyses for PSRL in FMDPs [Osband and Van Roy, 2014] and hierarchical posterior sampling [Hong et al., 2022b]. Our analysis shows that C-PSRL takes the best of both worlds by avoiding a direct dependence on the number of states in the regret (as in FMDPs) and without requiring a full causal graph prior (as in hierarchical posterior sampling). Moreover, we can analytically capture the dependency of the Bayesian regret on the number of causal edges known a priori and encoded in the (partial) causal graph prior. Finally, we empirically validate C-PSRL against two relevant baselines: PSRL with an uninformative prior, i.e., that does not model potential factorizations in the dynamics, and PSRL equipped with the full knowledge of the causal graph (an oracle prior). We carry out the comparison in simple yet illustrative domains, which show that exploiting a causal graph prior improves efficiency over uninformative priors while being only slightly inferior to the oracle prior.

In summary, the main contributions of this paper include the following:

- A novel problem formulation that links PSRL with a prior expressed as a partial causal graph to the problem of learning an FMDP with unknown factorization (Section 2);
- A method (C-PSRL) extending PSRL to exploit a partial causal graph prior (Section 3);
- The analysis of the Bayesian regret of C-PSRL, which is $\tilde{O}(\sqrt{K/2^\eta})$ where K is the total number of episodes and η is the degree of prior knowledge (Section 4);
- An ancillary result on causal discovery that shows how a (sparse) super-graph of the true causal graph can be extracted from a run of C-PSRL as a byproduct (Section 5);
- An experimental evaluation of the performance of C-PSRL against PSRL with uninformative or oracle priors in illustrative domains (Section 6).

The aim of this work is to enable the use of posterior sampling for RL in relevant applications through a causal perspective on prior specification. We believe this contribution can help to close the gap between PSRL research and actual adoption of PSRL in real-world problems.

---

[1]DTR is an example, where several causal relations affecting patient's conditions remain a mystery.

2 Problem formulation

In this section, we first provide preliminary background on graphical causal models (Section 2.1) and Markov decision processes (Section 2.2). Then, we explain how a causal graph on the variables of a Markov decision process induces a factorization of its dynamics (Section 2.3). Finally, we formalize the reinforcement learning problem in the presence of a causal graph prior (Section 2.4).

Notation. With few exceptions, we will denote a set or space as $\mathcal{A}$, their elements as $a \in \mathcal{A}$, constants or random variables with A, and functions as f. We denote $\Delta(\mathcal{A})$ the probability simplex over $\mathcal{A}$, and $[A]$ the set of integers $\{1, \ldots, A\}$. For a d-dimensional vector x, we define the *scope operator* $x[\mathcal{I}] := \bigotimes_{i \in \mathcal{I}} x_i$ for any set $\mathcal{I} \subseteq [d]$. When $\mathcal{I} = \{i\}$ is a singleton, we use $x[i]$ as a shortcut for $x[\{i\}]$. A recap of the notation, which is admittedly involved, can be found in Appendix A.

2.1 Causal graphs

Let $\mathcal{X} = \{X_j\}_{j=1}^{d_X}$ and $\mathcal{Y} = \{Y_j\}_{j=1}^{d_Y}$ be sets of random variables taking values $x_j, y_j \in [N]$ respectively, and let $p : \mathcal{X} \to \Delta(\mathcal{Y})$ a strictly positive probability density. Further, let $\mathcal{G} = (\mathcal{X}, \mathcal{Y}, z)$ be a bipartite Directed Acyclic Graph (DAG), or *bigraph*, having left variables $\mathcal{X}$, right variables $\mathcal{Y}$, and a set of edges $z \subseteq \mathcal{X} \times \mathcal{Y}$. We denote as z_j the parents of the variable $Y_j \in \mathcal{Y}$, such as $z_j = \{i \in [d_X] \mid (X_i, Y_j) \in z\}$ and $z = \bigcup_{j \in [d_Y]} \bigcup_{i \in z_j} \{(X_i, Y_j)\}$. We say that $\mathcal{G}$ is Z-sparse if $\max_{j \in [d_Y]} |z_j| \leq Z \leq d_X$. We call Z the *degree of sparseness* of $\mathcal{G}$.

The tuple $(p, \mathcal{G})$ is called a *graphical causal model* [Pearl, 2009] if p fulfills the Markov factorization property with respect to $\mathcal{G}$, that is $p(\mathcal{X}, \mathcal{Y}) = p(\mathcal{X})p(\mathcal{Y}|\mathcal{X}) = p(\mathcal{X}) \prod_{j \in [d_Y]} p_j(y[j]|x[z_j])$. Note that the causal model that we consider in this paper does not admit *confounding*. Further, we can exclude "vertical" edges in $\mathcal{Y} \times \mathcal{Y}$ and directed edges $\mathcal{Y} \times \mathcal{X}$, which allows to write $p(\mathcal{X}, \mathcal{Y}) = p(\mathcal{X})p(\mathcal{Y}|\mathcal{X})$. We call *causal graph* the component $\mathcal{G}$ of a graphical causal model.

2.2 Markov decision processes

A finite episodic Markov Decision Process [MDP, Puterman, 2014] is defined throug the tuple $\mathcal{M} := (\mathcal{S}, \mathcal{A}, p, r, \mu, H)$, where $\mathcal{S}$ is a state space of size S, $\mathcal{A}$ is an action space of size A, $p : \mathcal{S} \times \mathcal{A} \to \Delta(\mathcal{S})$ is a Markovian transition model such that $p(s'|s, a)$ denotes the conditional probability of the next state s' given the state s and action a, $r : \mathcal{S} \times \mathcal{A} \to \Delta([0, 1])$ is a reward function such that the reward collected performing action a in state s is distributed as $r(s, a)$ with mean $R(s, a) = \mathbb{E}[r(s, a)]$, $\mu \in \Delta(\mathcal{S})$ is the initial state distribution, $H < \infty$ is the episode horizon.

An agent interacts with the MDP as follows. First, the initial state is drawn $s_1 \sim \mu$. For each step $h < H$, the agent selects an action $a_h \in \mathcal{A}$. Then, they collect a reward $r_h \sim r(s_h, a_h)$ while the state transitions to $s_{h+1} \sim p(\cdot|s_h, a_h)$. The episode ends when s_H is reached.

The strategy from which the agent selects an action at each step is defined through a non-stationary, stochastic *policy* $\pi = \{\pi_h\}_{h \in [H]} \in \Pi$, where each $\pi_h : \mathcal{S} \to \Delta(\mathcal{A})$ is a function such that $\pi_h(a|s)$ denotes the conditional probability of selecting action a in state s at step h, and Π is the policy space. A policy $\pi \in \Pi$ can be evaluated through its value function $V_h^\pi : \mathcal{S} \to [0, H]$, which is the expected sum of rewards collected under π starting from state s at step h, i.e.,

$$V_h^\pi(s) := \mathbb{E}_\pi \left[\sum_{h'=h}^{H} R(s_{h'}, a_{h'}) \Big| s_h = s \right], \qquad \forall s \in \mathcal{S}, \ h \in [H].$$

We further define the value function of π in the MDP $\mathcal{M}$ under μ as $V_\mathcal{M}(\pi) := \mathbb{E}_{s \sim \mu}[V_1^\pi(s)]$.

Figure 1: **(Left)** Illustrative causal graph prior $\mathcal{G}_0$ with $d_X = 4, d_Y = 2$ features, degree of sparseness $Z = 3$. The hidden true graph $\mathcal{G}_{\mathcal{F}_*}$ includes all the edges in $\mathcal{G}_0$ plus the red-dashed edge $(3, 1)$. **(Right)** Visualization of $\mathcal{Z}$, the set of factorizations consistent with $\mathcal{G}_0$, which is the support of the hyper-prior P_0. The factorization z_* of the true FMDP $\mathcal{F}_*$ is highlighted in red.

2.3 Causal structure induces factorization

In the previous section, we formulated the MDP in a tabular representation, where each state (action) is identified by a unique symbol $s \in \mathcal{S}$ ($a \in \mathcal{A}$). However, in relevant real-world applications, the states and actions may be represented through a finite number of features, say d_S and d_A features respectively. The DTR problem is an example, where state features can be, e.g., blood pressure and glucose level, action features can be indicators on whether a particular medication is administered.

Let those state and action features be modeled by random variables in the interaction between an agent and the MDP, we can consider additional structure in the process by considering the causal graph of its variables, such that the value of a variable only depends on the values of its causal parents. Looking back to DTR, we might know that the value of the blood pressure at step $h+1$ only depends on its value at step h and whether a particular medication has been administered.

Formally, we can show that combining an MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, p, r, \mu, H)$ with a causal graph over its variables, which we denote as $\mathcal{G}_{\mathcal{M}} = (\mathcal{X}, \mathcal{Y}, z)$, gives a factored MDP [Boutilier et al., 2000]

$$\mathcal{F} = (\{\mathcal{X}_j\}_{j=1}^{d_X}, \{\mathcal{Y}_j, z_j, p_j, r_j\}_{j=1}^{d_Y}, \mu, H, Z, N),$$

where $\mathcal{X} = \mathcal{S} \times \mathcal{A} = \mathcal{X}_1 \times \ldots \times \mathcal{X}_{d_X}$ is a factored state-action space with $d_X = d_S + d_A$ discrete variables, $\mathcal{Y} = \mathcal{S} = \mathcal{Y}_1 \times \ldots \times \mathcal{Y}_{d_Y}$ is a factored state space with $d_Y = d_S$ variables, and z_j are the causal parents of each state variable, which are obtained from the edges z of $\mathcal{G}_{\mathcal{M}}$. Then, p is a *factored* transition model specified as $p(y|x) = \prod_{j=1}^{d_Y} p_j(y[j] \mid x[z_j]), \forall y \in \mathcal{Y}, x \in \mathcal{X}$, and r is a *factored* reward function $r(x) = \sum_{j=1}^{d_Y} r(x[z_j])$, with mean $R(x) = \sum_{j=1}^{d_Y} R(x[z_j]), \forall x \in \mathcal{X}$. Finally, $\mu \in \Delta(\mathcal{Y})$ and H are the initial state distribution and episode horizon as specified in $\mathcal{M}$, Z is the degree of sparseness of $\mathcal{G}_{\mathcal{M}}$, N is a constant such that all the variables are supported in $[N]$.

The interaction between an agent and the FMDP can be described exactly as we did in Section 2.2 for a tabular MDP, and the policies with their corresponding value functions are analogously defined. With the latter formalization of the FMDP induced by a causal graph, we now have all the components to introduce our learning problem in the next section.

2.4 Reinforcement learning with partial causal graph priors

In the previous section, we show how the prior knowledge of a causal graph over the MDP variables can be exploited to obtain an FMDP representation of the problem, which is well-known to allow for more efficient reinforcement learning thanks to the factorization of the transition model and reward function [Osband and Van Roy, 2014, Xu and Tewari, 2020, Tian et al., 2020, Chen et al., 2020, Talebi et al., 2021, Rosenberg and Mansour, 2021]. However, in several applications is unreasonable to assume prior knowledge of the full causal

graph, and causal identification is costly in general [Gillispie and Perlman, 2001, Shah and Peters, 2020]. Nonetheless, *some* prior knowledge of the causal graph, i.e., a portion of the edges, may be easily available. For instance, in a DTR problem some edges of the causal graph on patient's variables are commonly known, whereas several others are elusive.

In this paper, we study the reinforcement learning problem when a partial causal graph prior $\mathcal{G}_0 \subseteq \mathcal{G}_\mathcal{M}$ on the MDP $\mathcal{M}$ is available.[2] We formulate the learning problem in a Bayesian sense, in which the instance $\mathcal{F}_*$ is sampled from a prior distribution $P_{\mathcal{G}_0}$ *consistent* with the causal graph prior $\mathcal{G}_0$.[3] In Figure 1 (left), we illustrate both the causal graph prior $\mathcal{G}_0$ and the (hidden) true graph $\mathcal{G}_{\mathcal{F}_*}$ of the true instance $\mathcal{F}_*$. Analogously to previous works on Bayesian RL formulations, e.g., [Osband et al., 2013], we evaluate the performance of a learning algorithm in terms of its induced Bayesian regret.

Definition 1 (Bayesian Regret)**.** *Let $\mathfrak{A}$ a learning algorithm and let $P_{\mathcal{G}_0}$ a prior distribution on FMDPs consistent with the partial causal graph prior $\mathcal{G}_0$. The K-episodes Bayesian regret of $\mathfrak{A}$ is*

$$\mathcal{BR}(K) := \mathop{\mathbb{E}}_{\mathcal{F}_* \sim P_{\mathcal{G}_0}} \left[\sum_{k=1}^{K} V_*(\pi_*) - V_*(\pi_k) \right],$$

where $V_*(\pi) = V_{\mathcal{F}_*}(\pi)$ is the value of the policy π in $\mathcal{F}_*$ under μ, $\pi_* \in \arg\max_{\pi \in \Pi} V_*(\pi)$ is the optimal policy in $\mathcal{F}_*$, and π_k is the policy played by algorithm $\mathfrak{A}$ at step $k \in [K]$.

The Bayesian regret allows to evaluate a learning algorithm on average over multiple instances. This is particularly suitable in some domains, such as DTR, in which it is crucial to achieve a good performance of the treatment policy on different patients. In the next section, we introduce an algorithm that achieves a Bayesian regret rate that is sublinear in the number of episodes K.

3   Causal PSRL

To face the learning problem described in the previous section, we cannot naïvely apply the PSRL algorithm for FMDPs [Osband and Van Roy, 2014], since we cannot access the factorization z_* of the true instance $\mathcal{F}_*$, but only a causal graph prior $\mathcal{G}_0 = (\mathcal{X}, \mathcal{Y}, z_0)$ such that $z_0 \subseteq z_*$. Moreover, z_* is always *latent* in the interaction process, in which we can only observe state-action-reward realizations from $\mathcal{F}_*$. The latter can be consistent with several factorizations of the transition dynamics of $\mathcal{F}_*$, which means we can neither extract z_* directly from data. This is the common setting of *hierarchical Bayesian methods* [Hong et al., 2020, 2022a,b, Kveton et al., 2021], where a latent state is sampled from a latent hypothesis space on top of the hierarchy, which then conditions the sampling of the observed state down the hierarchy. In our setting, we can see the latent hypothesis space as the space of all the possible factorizations that are consistent with $\mathcal{G}_0$, whereas the observed states are the model parameters of the FMDP, from which we observe realizations. The algorithm that we propose, **C**ausal **PSRL** (C-PSRL), builds on this intuition to implement a principled hierarchical posterior sampling procedure to minimize the Bayesian regret exploiting the causal graph prior. We describe the details of the procedure below, whereas we report its pseudocode in Algorithm 1.

First, C-PSRL computes the set $\mathcal{Z}$, illustrated in Figure 1 (right), of the factorizations consistent with $\mathcal{G}_0$, i.e., which are both Z-sparse and include all of the edges in z_0 (line 2). Then, it specifies a parametric distribution P_0, called *hyper-prior*, over the latent hypothesis space $\mathcal{Z}$, and, for each $z \in \mathcal{Z}$, a further parametric distribution $P_0(\cdot|z)$, which is a *prior* on the model parameters, i.e., transition probabilities, conditioned on the latent state z (line 3).

---

[2]For two bigraphs $\mathcal{G}_\star = (\mathcal{X}, \mathcal{Y}, z_\star)$ and $\mathcal{G}_\bullet = (\mathcal{X}, \mathcal{Y}, z_\bullet)$, we let $\mathcal{G}_\star \subseteq \mathcal{G}_\bullet$ if $z_\star \subseteq z_\bullet$.

[3]We will specify in the next Section 3 how the prior $P_{\mathcal{G}_0}$ can be constructed.

Algorithm 1 Causal PSRL (C-PSRL)

1: **input**: causal graph prior $\mathcal{G}_0 \subseteq \mathcal{G}_{\mathcal{F}_*}$, degree of sparseness Z
2: Compute the set of consistent factorizations

$$\mathcal{Z} = \mathcal{Z}_1 \times \ldots \times \mathcal{Z}_{d_Y} = \left\{ z = \{z_j\}_{j \in [d_Y]} \ \middle| \ |z_j| < Z \text{ and } z_{0,j} \subseteq z_j \quad \forall j \in [d_Y] \right\}$$

3: Build the hyper-prior P_0 and the prior $P_0(\cdot|z)$ for each $z \in \mathcal{Z}$
4: **for** episode $k = 0, 1, \ldots, K-1$ **do**
5:     Sample $z \sim P_k$ and $p \sim P_k(\cdot|z)$ to build the FMDP $\mathcal{F}_k$
6:     Compute the policy $\pi_k \leftarrow \arg\max_{\pi \in \Pi} V_{\mathcal{F}_k}(\pi)$ collect an episode with π_k in $\mathcal{F}_*$
7:     Compute the posteriors P_{k+1} and $P_{k+1}(\cdot|z)$ with the collected data
8: **end for**

The former represents the agent's belief over the factorization of the true instance $\mathcal{F}_*$, the latter on the factored transition model p_*. [4]

Having translated the causal graph prior $\mathcal{G}_0$ into proper parametric prior distributions, C-PSRL executes a hierarchical posterior sampling procedure (lines 4-9). For each episode k, the algorithm sample a factorization z from the current hyper-prior P_k, and a transition model p from the prior $P_k(\cdot|z)$, such that p is factored according to z (line 5). With these two objects, it builds the FMDP $\mathcal{F}_k$ (line 5), for which it computes the optimal policy π_k solving the corresponding planning problem (line 6). Finally, the policy π_k is deployed on the true instance $\mathcal{F}_*$ for one episode (line 7) to collect the evidence that serves to compute the posterior updates of the prior and hyper-prior (line 8).

As we shall see, the described Algorithm 1 has compelling statistical properties, as it suffers a sublinear regret in the number of episodes K (Section 4) while providing a notion of causal discovery as a byproduct (Section 5), and showcases promising empirical performance (Section 6).

Recipe. Three key ingredients concur to make the algorithm successful. First and foremost, C-PSRL links RL of an FMDP $\mathcal{F}_*$ with unknown factorization to a hierarchical Bayesian learning problem, in which the factorization acts as a latent state on top of the hierarchy, and the transition probabilities are the observed state down the hierarchy. Secondly, C-PSRL exploits the causal graph prior $\mathcal{G}_0$ to reduce the size of the latent hypothesis space $\mathcal{Z}$, which is super-exponential in the number of features d_X, d_Y of $\mathcal{F}_*$ in general [Robinson, 1973]. Finally, C-PSRL harnesses the specific causal structure of the problem to get a factorization z (line 5) through independent sampling of the parents $z_j \in \mathcal{Z}_j$ for each Y_j, which significantly reduces the number of hyper-prior parameters. Crucially, this can be done as we do not admit "vertical" edges in $\mathcal{Y}$ and edges directed from $\mathcal{Y}$ to $\mathcal{X}$ in our hypothesis space, such that it is impossible to select parents' assignment that leads to a cycle.

Degree of sparseness. C-PSRL takes as input (line 1) the degree of sparseness Z of the true FMDP $\mathcal{F}_*$, which might be unknown in practice. In that case, Z can be seen as an hyper-parameter of the algorithm, which can be either implied through domain expertise or tuned independently.

Planning in FMDPs. C-PSRL requires exact planning in a FMDP (line 11), which is intractable in general [Mundhenk et al., 2000, Lusena et al., 2001]. While we do not directly address computational issues in this paper, we note that efficient approximation schemes have been developed for this problem [Guestrin et al., 2003]. Moreover, under linear realizability assumptions for the transition model or value functions, i.e., they can be represented through a linear combination of known features, exact planning methods exist [Yang and Wang, 2019, Jin et al., 2020b, Deng et al., 2022].

---

[4]A description of parametric distributions P_0 and $P_0(\cdot|z)$ and their posterior updates is in Appendix B.

4   Regret analysis of C-PSRL

In this section, we study the Bayesian regret induced by C-PSRL with a Z-sparse causal graph prior $\mathcal{G}_0 = (\mathcal{X}, \mathcal{Y}, z_0)$. First, we define the *degree of prior knowledge* $\eta \leq \min_{j \in [d_Y]} |z_{0,j}|$, which is a lower bound on the number of causal parents revealed by the prior $\mathcal{G}_0$ for each state variable Y_j of the true instance $\mathcal{F}_*$. With this definition, we provide an upper bound on the Bayesian regret incurred by C-PSRL, which we then discuss in Section 4.1.[5]

Theorem 4.1. *Let $\mathcal{G}_0$ be a causal graph prior with degree of sparseness Z and degree of prior knowledge η. The K-episodes Bayesian regret incurred by C-PSRL is*

$$\mathcal{BR}(K) = \tilde{O}\left(\left(H^{5/2}N^{1+Z/2}d_Y + \sqrt{H2^{d_X-\eta}}\right)\sqrt{K}\right).$$

While we defer the proof of the result to Appendix D, we report a sketch below.

Step 1.   The first step of our proof bridges the previous analyses of a hierarchical version of PSRL, which is reported in [Hong et al., 2022b], with the one of PSRL for factored MDPs [Osband and Van Roy, 2014]. In short, we can decompose the Bayesian regret (see Definition 1) as

$$\mathcal{BR}(K) = \mathbb{E}\left[\sum_{k=1}^{K}\mathbb{E}_k\left[V_*(\pi_*) - \overline{V}_k(\pi_*, Z_*)\right]\right] + \mathbb{E}\left[\sum_{k=1}^{K}\mathbb{E}_k\left[\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\right]\right]$$

where $\mathbb{E}_k[\cdot]$ is the conditional expectation given the evidence collected until episode k, and $\overline{V}_k(\pi, z) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot|z)}[V_{\mathcal{F}}(\pi)]$ is the value function of π on average over the posterior $P_k(\cdot|z)$. Informally, the first term captures the regret due to the concentration of the *posterior* $P_k(\cdot|z_*)$ around the true transition model p_* having fixed the true factorization z_*. Instead, the second term captures the regret due to the concentration of the *hyper-posterior* P_k around the true factorization z_*. Through a non-trivial adaptation of the analysis in [Hong et al., 2022b] to the FMDP setting, we can bound each term separately to obtain $\tilde{O}((H^{5/2}N^{1+Z/2}d_Y + \sqrt{H|\mathcal{Z}|})\sqrt{K})$.

Step 2.   The upper bound of the previous step is close to the final result up to a factor $\sqrt{|\mathcal{Z}|}$ related the size of the latent hypothesis space. Since C-PSRL performs local sampling from the product space $\mathcal{Z} = \mathcal{Z}_1 \times \ldots \times \mathcal{Z}_{d_Y}$, by combining independent samples $z_j \in \mathcal{Z}_j$ for each variable Y_j as we briefly explained in Section 3, we can refine the dependence in $|\mathcal{Z}|$ to $\max_{j \in [d_Y]}|\mathcal{Z}_j| \leq |\mathcal{Z}|$.

Step 3.   Finally, to obtain the final rate reported in Theorem 4.1, we have to capture the dependency in the degree of prior knowledge η in the Bayesian regret by upper bounding $\max_{j \in [d_Y]}|\mathcal{Z}_j|$ as

$$\max_{j \in [d_Y]}|\mathcal{Z}_j| = \sum_{i=0}^{Z-\eta}\binom{d_X - \eta}{i} \leq 2^{d_X-\eta}.$$

4.1   Discussion of the Bayesian regret

The regret bound in Theorem 4.1 contains two terms, which informally capture the regret to learn the transition model having the true factorization (left), and to learn the true factorization (right).

The first term is typical in previous analyses of vanilla posterior sampling. Especially, the best known rate for the MDP setting is $\tilde{O}(H\sqrt{SAK})$ [Osband and Van Roy, 2017]. In a FMDP setting with known factorization, the direct dependencies with the size S, A of the state and action spaces can be refined to obtain $\tilde{O}(Hd_Y^{3/2}N^{Z/2}\sqrt{K})$ [Osband and Van Roy,

---

[5]We report the regret rate with the common "Big-O" notation, in which $\tilde{O}$ hides logarithmic factors.

2014]. Our rate includes additional factors of H and N, but a better dependency on the number of state features d_Y.

The second term of the regret rate is instead unique to hierarchical Bayesian settings, which include an additional source of randomization in the sampling of the latent state from the hyper-prior. In Theorem 4.1, we are able to express this term in the degree of prior knowledge η, resulting in a rate $\tilde{O}(\sqrt{K/2^\eta})$. The latter naturally demonstrates that a richer causal graph prior $\mathcal{G}_0$ will benefit the efficiency of PSRL, bringing the regret rate closer to the one for an FMDP with known factorization.

We believe the rate in Theorem 4.1 is shedding light on how prior causal knowledge, here expressed through a partial causal graph, impacts on the efficiency of posterior sampling for RL.

5 C-PSRL embeds a notion of causal discovery

In this section, we provide an ancillary result that links Bayesian regret minimization with C-PSRL to a notion of causal discovery, which we call *weak causal discovery*. Especially, we show that we can extract a Z-sparse super-graph of the causal graph $\mathcal{G}_{\mathcal{F}_*}$ of the true instance $\mathcal{F}_*$ as a byproduct.

A run of C-PSRL produces a sequence $\{\pi_k\}_{k=0}^{K-1}$ of optimal policies for the FMDPs $\{\mathcal{F}_k\}_{k=0}^{K-1}$ drawn from the posteriors. Every FMDP $\mathcal{F}_k$ is linked to a corresponding graph (or factorization) $\mathcal{G}_{\mathcal{F}_k} = (\mathcal{X}, \mathcal{Y}, z_k)$, where $z_k \sim P_k$ is sampled from the hyper-posterior. Note that the algorithm does not enforce any causal meaning to the edges z_k of $\mathcal{G}_{\mathcal{F}_k}$. Nonetheless, we aim to show that we can extract a Z-sparse super-graph of $\mathcal{G}_{\mathcal{F}_*}$ from the sequence $\{\mathcal{F}_k\}_{k=0}^{K-1}$ with high probability.

First, we need to assume that any misspecification in $\mathcal{G}_{\mathcal{F}_k}$ negatively affects the value function of π_k. Thus, we extend the traditional notion of causal minimality [Spirtes et al., 2000] to value functions.

Definition 2 (ϵ-Value Minimality). *An FMDP $\mathcal{F}$ fulfills ϵ-value minimality, if for any FMDP $\mathcal{F}'$ encoding a proper subgraph of $\mathcal{G}_{\mathcal{F}}$, i.e., $\mathcal{G}_{\mathcal{F}'} \subset \mathcal{G}_{\mathcal{F}}$, it holds that $V_{\mathcal{F}}^* > V_{\mathcal{F}'}^* + \epsilon$, where $V_{\mathcal{F}}^*, V_{\mathcal{F}'}^*$ are the value functions of the optimal policies in $\mathcal{F}, \mathcal{F}'$ respectively.*

Then, as a corollary of Theorem 4.1, we can prove the following result.

Corollary 5.1 (Weak Causal Discovery). *Let $\mathcal{F}_*$ be an FMDP in which the transition model p_* fulfills the causal minimality assumption with respect to $\mathcal{G}_{\mathcal{F}_*}$, and let $\mathcal{F}_*$ fulfill ϵ-value minimality. Then, $\mathcal{G}_{\mathcal{F}_*} \subseteq \mathcal{G}_{\mathcal{F}_K}$ holds with high probability, where $\mathcal{G}_{\mathcal{F}_K}$ is a Z-sparse graph randomly selected within the sequence $\{\mathcal{G}_{\mathcal{F}_k}\}_{k=0}^{K-1}$ produced by C-PSRL over $K = \tilde{O}(H^5 d_Y^2 2^{d_X - \eta}/\epsilon^2)$ episodes.*

The latter result shows that C-PSRL discovers the causal relationships between the FMDP variables, but cannot easily prune the non-causal edges, making $\mathcal{G}_{\mathcal{F}_K}$ a super-graph of $\mathcal{G}_{\mathcal{F}_*}$. In Appendix C, we report a detailed derivation of the previous result. Interestingly, Corollary 5.1 suggests a direct link between regret minimization in a FMDP with unknown factorization and a (weak) notion of causal discovery, which might be further explored in future works.

6 Experiments

In this section, we provide experiments to both support the design of C-PSRL (Section 3) and validate its regret rate (Section 4). We consider two simple yet illustrative domains. The first, which we call *Random FMDP*, benchmarks the performance of C-PSRL on randomly generated FMDP instances, a setting akin to the Bayesian learning problem

Figure 2: **(a,b)** Regret and model error as a function of the episodes in the Random FMDP domain with $d_X = 9, d_Y = 6, Z = 5, N = 2, H = 100$. **(c,d)** Regret as a function of the episodes in Taxi 3×3 with $d_X = 5, d_Y = 4, Z = 5, N = [3, 3, 2, 1, 6], H = 10$, Taxi 5×5 with $d_X = 5, d_Y = 4, Z = 5, N = [5, 5, 2, 1, 6], H = 15$. The plots report the mean and 95% c.i. over 20 runs.

(see Section 2.4) that we considered in previous sections. The latter is a traditional *Taxi* environment [Dietterich, 2000], which is naturally factored and hints at a potential application. In those domains, we compare C-PSRL against two natural baselines: PSRL for tabular MDPs [Strens, 2000], and Factored PSRL (F-PSRL), which extends PSRL to factored MDP settings [Osband and Van Roy, 2014]. Note that F-PSRL is equivalent to an instance of C-PSRL that receives the true causal graph prior as input, i.e., has an oracle prior.

Random FMDPs. An FMDP (relevant parameters are reported in the caption of Figure 2) is sampled uniformly from the prior specified through a random causal graph, which is Z-sparse with at least two edges for every state variable ($\eta = 2$). Then, the regret is minimized by running PSRL, F-PSRL, and C-PSRL ($\eta = 2$) for 500 episodes. Figure 2a shows that C-PSRL achieves a regret that is significantly smaller than PSRL, thus outperforming the baseline with an uninformative prior, while being surprisingly close to F-PSRL, having the oracle prior. Indeed, C-PSRL resulted efficient in estimating the transition model of the sampled FMDP, as we can see from Figure 2b, which reports the ℓ_1 distance between the true model p_* and the p_k sampled by the algorithm at episode k.

Taxi. For the Taxi domain, we use the common Gym implementation [Brockman et al., 2016]. In this environment, a taxi driver needs to pick up a passenger at a specific location, and then it has to bring the passenger to their destination. The environment is represented as a grid, with some special cells identifying the passenger location and destination. As reported in Simão and Spaan [2019], this domain is inherently factored since the state space is represented by four independent features: The position of the taxi (row and column), the passenger's location and whether they are on the taxi, and the destination. We perform the experiment on two grids with varying size (3×3 and 5×5 respectively), for which we report the relevant parameters in Figure 2. Here we compare the proposed algorithm C-PSRL ($\eta = 2$) with PSRL, while F-PSRL is omitted as the knowledge of the oracle prior is not available. Both algorithms converge to a good policy eventually in the smaller grid (see the regret in Figure 2c). Instead, when the size of the grid increases, PSRL is still suffering a linear regret after 400 episodes, whereas C-PSRL succeeds in finding a good policy efficiently (see Figure 2d). Notably, this domain resembles the problem of learning optimal routing in a taxi service, and our results show that exploiting common knowledge (such as that the location of the taxi and passenger's destination) in the form of a causal graph prior can be a game changer.

7 Related work

We revise here the most relevant related work in posterior sampling, FMDPs, and causal RL.

Posterior sampling. Thompson sampling [Thompson, 1933] is a well-known Bayesian algorithm that has been extensively analyzed in both multi-armed bandit problems and

reinforcement learning [Kaufmann et al., 2012, Agrawal and Goyal, 2012, Osband et al., 2013, Osband and Van Roy, 2017]. Specifically, Osband and Van Roy [2017] provides a regret rate $\tilde{O}(H\sqrt{SAK})$ for vanilla Thompson sampling in RL, which is called the PSRL algorithm. Recently, other works adapted Thompson sampling to hierarchical Bayesian problems [Hong et al., 2020, 2022a,b, Kveton et al., 2021]. Mixture Thompson Sampling (MixTS) [Hong et al., 2022b], which is similar to PSRL but samples the unknown MDP from a mixture prior, is arguably the closest to our setting. In this paper, we take inspiration from their algorithm to design C-PSRL and derive its analysis, even though, instead of their tabular setting, we tackle a fundamentally different problem on factored MDPs resulting from a casual graph prior, which induces unique challenges.

Factored MDPs.   Our setting is also related to FMDPs [Boutilier et al., 2000]. Previous works considered reinforcement learning in FMDPs with either known [Osband and Van Roy, 2014, Xu and Tewari, 2020, Talebi et al., 2021, Tian et al., 2020, Chen et al., 2020] or unknown [Strehl et al., 2007, Vigorito and Barto, 2009, Diuk et al., 2009, Chakraborty and Stone, 2011, Hallak et al., 2015, Guo and Brunskill, 2017, Rosenberg and Mansour, 2021] factorization. The PSRL algorithm has been adapted to both finite-horizon [Osband and Van Roy, 2014] and infinite-horizon [Xu and Tewari, 2020] FMDPs. The former assumes knowledge of the factorization, which is close to our setting with an oracle prior, and provides Bayesian regret rate of order $\tilde{O}(Hd_Y^{3/2}N^{Z/2}\sqrt{K})$. Previous works have also studied reinforcement learning in FMDPs in a frequentist sense, either with known [Chen et al., 2020] or unknown [Rosenberg and Mansour, 2021] factorization. Rosenberg and Mansour [2021] employ an optimistic method that is orthogonal to ours, whereas they leave as an open problem capturing the effect of prior knowledge, for which we provide answers in a Bayesian setting.

Causal RL.   Previous works have addressed RL with a causal perspective (see [Kaddour et al., 2022, Chapter 7] for a survey of such methods). Those works typically exploit causal principles to obtain compact representations of states and transitions [Tomar et al., 2021, Gasse et al., 2021], or to pursue generalization across tasks and environments [Zhang et al., 2020, Huang et al., 2022, Feng et al., 2022, Mutti et al., 2023]. Closer to our setting, Lu et al. [2022] address the problem of exploiting prior causal knowledge to learn in both MDPs and FMDPs. Our work differs from theirs in two key aspects: We show how to exploit a partial causal graph prior instead of assuming knowledge of the full causal graph, and we consider a Bayesian formulation of the problem while they tackle a frequentist setting through optimism principles. Finally, the work in [Zhang, 2020b] shows an interesting application of causal RL for designing treatments in a DTR problem.

8   Conclusion

In this paper, we presented how to exploit prior knowledge expressed through a partial causal graph to improve the statistical efficiency of reinforcement learning. Before reporting some concluding remarks, it is worth commenting on where such a causal graph prior might be originated from.

Exploiting experts' knowledge.   One natural application of our methodology is to exploit domain-specific knowledge coming from experts. In several domains, e.g., medical or scientific applications, expert practitioners have some knowledge over the causal relationships between the domain's variables. However, they might not have a full picture of the causal structure, especially when they face complex systems such as the human body or biological processes. Our methodology allows those practitioners to easily encode their partial knowledge into a graph prior, instead of having to deal with technically involved Bayesian statistics to specify parametric prior distributions, and then let C-PSRL figure out a competent decision policy with the given information.

Exploiting causal discovery. Identifying the causal graph over domain's variables, which is usually referred as causal discovery, is a main focus of *causality* [Pearl, 2009, Chapter 3]. The literature provides plenty of methods to perform causal discovery from data [Peters et al., 2017, Chapter 4], including learning causal variables and their relationships in MDP settings [Zhang et al., 2020, Mutti et al., 2023]. However, learning the full causal graph, even when it is represented with a bigraph as in MDP settings [Mutti et al., 2023], can be statistically costly or even prohibitive [Gillispie and Perlman, 2001, Wadhwa and Dong, 2021]. Moreover, not all the causal edges are guaranteed to transfer across environments [Mutti et al., 2023], which would force to perform causal discovery anew for any slight variation of the domain (e.g., changing the patient in a DTR setting). Our methodology allows to focus on learning the *universal* causal relationships [Mutti et al., 2023], which transfer across environments, e.g., different patients, and then specify the prior through a partial causal graph.

The latter paragraphs describe two scenarios in which our work enhance the applicability of PSRL, bridging the gap between how the prior might be specified in practical applications and what previous methods currently require, i.e., a parametric prior distribution. To summarize our contributions, we first provided a Bayesian formulation of reinforcement learning with prior knowledge expressed through a partial causal graph. Then, we presented an algorithm, C-PSRL, tailored for the latter problem, and we analyzed its regret to obtain a rate that is sublinear in the number of episodes and shows a direct dependence with the degree of causal knowledge. Finally, we derived an ancillary result to show that C-PSRL embeds a notion of causal discovery, and we provided an empirical validation of the algorithm against relevant baselines. C-PSRL resulted nearly competitive with F-PSRL, which enjoys a richer prior, while clearly outperforming PSRL with an uninformative prior.

Future works may derive a tighter analysis of the Bayesian regret of C-PSRL, as well as a stronger causal discovery result that allows to recover a minimal causal graph instead of a super-graph. Finally, another important future direction is to address computational issues inherent to planning in FMDPs to scale the implementation of C-PSRL to complex domains.

References

S. Agrawal and N. Goyal. Analysis of Thompson sampling for the multi-armed bandit problem. In *Conference on Learning Theory*, 2012.

R. Bellman. Dynamic programming. *Princeton University Press*, 1957.

C. Boutilier, R. Dearden, and M. Goldszmidt. Stochastic dynamic programming with factored representations. *Artificial Intelligence*, 121(1-2):49–107, 2000.

G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, and W. Zaremba. Openai gym. *arXiv preprint arXiv:1606.01540*, 2016.

D. Chakraborty and P. Stone. Structure learning in ergodic factored mdps without knowledge of the transition function's in-degree. In *International Conference on Machine Learning*, 2011.

O. Chapelle and L. Li. An empirical evaluation of Thompson sampling. In *Advances in Neural Information Processing Systems*, 2011.

X. Chen, J. Hu, L. Li, and L. Wang. Efficient reinforcement learning in factored mdps with application to constrained rl. In *International Conference on Learning Representations*, 2020.

Z. Deng, S. Devic, and B. Juba. Polynomial time reinforcement learning in factored state mdps with linear value functions. In *International Conference on Artificial Intelligence and Statistics*, 2022.

T. G. Dietterich. Hierarchical reinforcement learning with the maxq value function decomposition. *Journal of Artificial Intelligence Research*, 13:227–303, 2000.

C. Diuk, L. Li, and B. R. Leffler. The adaptive k-meteorologists problem and its application to structure learning and feature selection in reinforcement learning. In *International Conference on Machine Learning*, 2009.

F. Feng, B. Huang, K. Zhang, and S. Magliacane. Factored adaptation for non-stationary reinforcement learning. In *Advances in Neural Information Processing Systems*, 2022. doi: 10.48550/ARXIV.2203.16582.

M. Gasse, D. Grasset, G. Gaudron, and P.-Y. Oudeyer. Causal reinforcement learning using observational and interventional data. *arXiv preprint arXiv:2106.14421*, 2021.

S. B. Gillispie and M. D. Perlman. Enumerating Markov equivalence classes of acyclic digraph models. In *Uncertainty in Artificial Intelligence*, 2001.

C. Guestrin, D. Koller, R. Parr, and S. Venkataraman. Efficient solution algorithms for factored mdps. *Journal of Artificial Intelligence Research*, 19:399–468, 2003.

Z. D. Guo and E. Brunskill. Sample efficient feature selection for factored mdps. *arXiv preprint arXiv:1703.03454*, 2017.

A. Hallak, F. Schnitzler, T. Mann, and S. Mannor. Off-policy model-based learning under unknown factored dynamics. In *International Conference on Machine Learning*, 2015.

J. Hong, B. Kveton, M. Zaheer, Y. Chow, A. Ahmed, and C. Boutilier. Latent bandits revisited. In *Advances in Neural Information Processing Systems*, 2020.

J. Hong, B. Kveton, M. Zaheer, and M. Ghavamzadeh. Hierarchical bayesian bandits. In *International Conference on Artificial Intelligence and Statistics*, 2022a.

J. Hong, B. Kveton, M. Zaheer, M. Ghavamzadeh, and C. Boutilier. Thompson sampling with a mixture prior. In *International Conference on Artificial Intelligence and Statistics*, 2022b.

B. Huang, F. Feng, C. Lu, S. Magliacane, and K. Zhang. Adarl: What, where, and how to adapt in transfer reinforcement learning. In *International Conference on Learning Representations*, 2022.

C. Jin, Z. Allen-Zhu, S. Bubeck, and M. I. Jordan. Is Q-learning provably efficient? In *Advances in Neural Information Processing Systems*, 2018.

C. Jin, A. Krishnamurthy, M. Simchowitz, and T. Yu. Reward-free exploration for reinforcement learning. In *International Conference on Machine Learning*, 2020a.

C. Jin, Z. Yang, Z. Wang, and M. I. Jordan. Provably efficient reinforcement learning with linear function approximation. In *Conference on Learning Theory*, 2020b.

J. Kaddour, A. Lynch, Q. Liu, M. J. Kusner, and R. Silva. Causal machine learning: A survey and open problems. *arXiv preprint arXiv:2206.15475*, 2022.

E. Kaufmann, N. Korda, and R. Munos. Thompson sampling: An asymptotically optimal finite-time analysis. In *Algorithmic Learning Theory*, 2012.

B. Kveton, C. wei Hsu, C. Boutilier, C. Szepesvari, M. Zaheer, M. Mladenov, and M. Konobeev. Meta-Thompson sampling. In *International Conference on Machine Learning*, 2021.

T. Lattimore and C. Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2020.

Y. Lu, A. Meisami, and A. Tewari. Efficient reinforcement learning with prior causal knowledge. In *Conference on Causal Learning and Reasoning*, 2022.

C. Lusena, J. Goldsmith, and M. Mundhenk. Nonapproximability results for partially observable markov decision processes. *Journal of Artificial Intelligence Research*, 14: 83–103, 2001.

O. Marchal and J. Arbel. On the sub-Gaussianity of the Beta and Dirichlet distributions. *Electronic Communications in Probability*, 22:1–14, 2017.

A. Marx, A. Gretton, and J. M. Mooij. A weaker faithfulness assumption based on triple interactions. In *Uncertainty in Artificial Intelligence*, 2021.

M. Mundhenk, J. Goldsmith, C. Lusena, and E. Allender. Complexity of finite-horizon Markov decision process problems. *Journal of the ACM (JACM)*, 47(4):681–720, 2000.

S. A. Murphy. Optimal dynamic treatment regimes. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 65(2):331–355, 2003.

M. Mutti, R. De Santi, E. Rossi, J. F. Calderon, M. Bronstein, and M. Restelli. Provably efficient causal model-based reinforcement learning for systematic generalization. In *AAAI Conference on Artificial Intelligence*, 2023.

I. Osband and B. Van Roy. Near-optimal reinforcement learning in factored mdps. In *Advances in Neural Information Processing Systems*, 2014.

I. Osband and B. Van Roy. Why is posterior sampling better than optimism for reinforcement learning? In *International Conference on Machine Learning*, 2017.

I. Osband, D. Russo, and B. Van Roy. (More) efficient reinforcement learning via posterior sampling. In *Advances in Neural Information Processing Systems*, 2013.

J. Pearl. *Causality*. Cambridge University Press, 2009.

J. Peters, D. Janzing, and B. Schölkopf. *Elements of causal inference: Foundations and learning algorithms*. The MIT Press, 2017.

M. L. Puterman. *Markov decision processes: Discrete stochastic dynamic programming*. John Wiley & Sons, 2014.

R. W. Robinson. Counting labeled acyclic digraphs. *New Directions in the Theory of Graphs*, pages 239–273, 1973.

A. Rosenberg and Y. Mansour. Oracle-efficient regret minimization in factored mdps with unknown structure. In *Advances in Neural Information Processing Systems*, 2021.

D. Russo and B. Van Roy. Learning to optimize via posterior sampling. *Mathematics of Operations Research*, 39(4):1221–1243, 2014.

R. Shah and J. Peters. The hardness of conditional independence testing and the generalised covariance measure. *Annals of Statistics*, 48(3):1514–1538, 2020.

T. D. Simão and M. T. Spaan. Safe policy improvement with baseline bootstrapping in factored environments. In *AAAI Conference on Artificial Intelligence*, 2019.

P. Spirtes, C. N. Glymour, R. Scheines, and D. Heckerman. *Causation, prediction, and search*. The MIT Press, 2000.

A. L. Strehl, C. Diuk, and M. L. Littman. Efficient structure learning in factored-state mdps. In *AAAI Conference on Artificial Intelligence*, 2007.

M. Strens. A Bayesian framework for reinforcement learning. In *International Conference on Machine Learning*, 2000.

R. S. Sutton and A. G. Barto. *Reinforcement learning: An introduction*. MIT press, 2018.

M. S. Talebi, A. Jonsson, and O. Maillard. Improved exploration in factored average-reward mdps. In *International Conference on Artificial Intelligence and Statistics*, 2021.

W. R. Thompson. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 1933.

Y. Tian, J. Qian, and S. Sra. Towards minimax optimal reinforcement learning in factored Markov decision processes. In *Advances in Neural Information Processing Systems*, 2020.

M. Tomar, A. Zhang, R. Calandra, M. E. Taylor, and J. Pineau. Model-invariant state abstractions for model-based reinforcement learning. *arXiv preprint arXiv:2102.09850*, 2021.

C. M. Vigorito and A. G. Barto. Incremental structure learning in factored mdps with continuous states and actions. *University of Massachusetts Amherst-Department of Computer Science, Technical Report*, 2009.

S. Wadhwa and R. Dong. On the sample complexity of causal discovery and the value of domain expertise. *arXiv preprint arXiv:2102.03274*, 2021.

Z. Xu and A. Tewari. Reinforcement learning in factored mdps: Oracle-efficient algorithms and tighter regret bounds for the non-episodic setting. In *Advances in Neural Information Processing Systems*, 2020.

L. Yang and M. Wang. Sample-optimal parametric Q-learning using linearly additive features. In *International Conference on Machine Learning*, 2019.

A. Zhang, C. Lyle, S. Sodhani, A. Filos, M. Kwiatkowska, J. Pineau, Y. Gal, and D. Precup. Invariant causal prediction for block mdps. In *International Conference on Machine Learning*, 2020.

J. Zhang. A comparison of three Occam's razors for Markovian causal models. *The British Journal for the Philosophy of Science*, 2020a.

J. Zhang. Designing optimal dynamic treatment regimes: A causal reinforcement learning approach. In *International Conference on Machine Learning*, 2020b.

J. Zhang and P. Spirtes. Detection of unfaithfulness and robust causal inference. *Minds and Machines*, 18(2):239–271, 2008.

Appendix

A   List of symbols

<div align="center">Basic mathematical objects</div>

$\mathcal{A}$	$\triangleq$	Set or space
A	$\triangleq$	Constant or random variable
a	$\triangleq$	Element of a set
$\Delta(\mathcal{A})$	$\triangleq$	Probability simplex over $\mathcal{A}$
$f : \mathcal{A} \to \mathcal{B}$	$\triangleq$	Function from $\mathcal{A}$ to $\mathcal{B}$
$[A]$	$\triangleq$	Set of integers $[A] = \{1, \ldots, A\}$
$x[\mathcal{I}]$	$\triangleq$	Scope operator $x[\mathcal{I}] := \bigotimes_{i \in \mathcal{I}} x_i$ for any set $\mathcal{I} \subseteq [d]$, $x \in \mathbb{R}^d$

<div align="center">Causal graph</div>

$\mathcal{G}$	$\triangleq$	Directed acyclic bigraph $\mathcal{G} = (\mathcal{X}, \mathcal{Y}, z)$
$\mathcal{X}$	$\triangleq$	Set of d_X random variables $\{X_j\}_{j=1}^{d_X}$ taking values $x_j \in [N]$
$\mathcal{Y}$	$\triangleq$	Set of d_Y random variables $\{Y_j\}_{j=1}^{d_Y}$ taking values $y_j \in [N]$
z	$\triangleq$	Directed edges $z \subseteq \mathcal{X} \times \mathcal{Y}$
z_j	$\triangleq$	Parents of Y_j such that $z_j = \{i \mid (X_i, Y_j) \in z\}$
Z	$\triangleq$	Degree of sparseness such that $|z_j| < Z, \forall j \in [d_Y]$
N	$\triangleq$	Size of the support of random variables

<div align="center">MDP</div>

$\mathcal{M}$	$\triangleq$	Markov decision process $\mathcal{M} = (\mathcal{S}, \mathcal{A}, p, r, \mu, H)$
$\mathcal{S}$	$\triangleq$	State space
$\mathcal{A}$	$\triangleq$	Action space
p	$\triangleq$	Transition model $p : \mathcal{S} \times \mathcal{A} \to \Delta(\mathcal{S})$
r	$\triangleq$	Reward function $r : \mathcal{S} \times \mathcal{A} \to \Delta([0, 1])$
μ	$\triangleq$	Initial state distribution $\mu \in \Delta(\mathcal{S})$
H	$\triangleq$	Episode horizon $H < \infty$
S	$\triangleq$	Size of the state space $S = |\mathcal{S}|$
A	$\triangleq$	Size of the action space $A = |\mathcal{A}|$
s	$\triangleq$	State $s \in \mathcal{S}$
a	$\triangleq$	Action $a \in \mathcal{A}$
$R(s, a)$	$\triangleq$	Mean reward $\mathbb{E}[r(s, a)]$

<div align="center">Factored MDP</div>

$\mathcal{F}$	$\triangleq$	Factored Markov Decision Process $\mathcal{F} = (\{\mathcal{X}_j\}_{j=1}^{d_X}, \{\mathcal{Y}_j, z_j, p_j, r_j, \}_{j=1}^{d_Y}, \mu, H, Z, N)$
d_X	$\triangleq$	Number of state-action variables
d_Y	$\triangleq$	Number of state variables
$\mathcal{X}$	$\triangleq$	Factored state-action space $\mathcal{X} = \mathcal{X}_1 \times \ldots \times \mathcal{X}_{d_X}$
$\mathcal{Y}$	$\triangleq$	Factored state space $\mathcal{Y} = \mathcal{Y}_1 \times \ldots \times \mathcal{Y}_{d_Y}$
z	$\triangleq$	Directed edges $z \subseteq \mathcal{X} \times \mathcal{Y}$, i.e., a *factorization*
z_j	$\triangleq$	Parents of Y_j such that $z_j = \{i \mid (X_i, Y_j) \in z\}$
p	$\triangleq$	Factored transition model $p(y|x) = \prod_{j=1}^{d_Y} p_j(y[j] \mid x[z_j])$
r	$\triangleq$	Factored reward function $r(x) = \sum_{j=1}^{d_Y} r_j(x[z_j])$
μ	$\triangleq$	Initial state distribution $\mu \in \Delta(\mathcal{Y})$
H	$\triangleq$	Episode horizon $H < \infty$
Z	$\triangleq$	Degree of sparseness such that $|z_j| < Z, \forall j \in [d_Y]$
N	$\triangleq$	Size of the support of state and action variables

<div align="center">Learning problem</div>

K	$\triangleq$	Number of episodes

k	$\triangleq$	Episode index
h	$\triangleq$	Step index
$\mathcal{Z}$	$\triangleq$	Space of the consistent factorizations $\mathcal{Z} \subseteq \mathcal{X} \times \mathcal{Y}$
$\mathcal{Z}_j$	$\triangleq$	Space of the consistent parents of Y_j such that $\mathcal{Z} = \mathcal{Z}_1 \times \ldots \mathcal{Z}_{d_Y}$
P_0	$\triangleq$	Hyper-prior on the factorizations consistent with $\mathcal{G}_0$ (supported in $\mathcal{Z}$)
P_k	$\triangleq$	Posterior of the hyper-prior P_0 at episode $k \in [K]$
$P_0(\cdot|z)$	$\triangleq$	Prior on the FMDPs with factorization z
$P_k(\cdot|z)$	$\triangleq$	Posterior of the prior $P_0(\cdot|z)$ at episode $k \in [K]$
$P_{\mathcal{G}_0}$	$\triangleq$	Prior on the FMDPs consistent with $\mathcal{G}_0$ such that
		$$P_{\mathcal{G}_0}(\mathcal{F}) = P_0(p_{\mathcal{F}}|z_{\mathcal{F}})P_0(z_{\mathcal{F}})$$
$\mathcal{BR}(K)$	$\triangleq$	K-episodes Bayesian regret

Regret analysis

Ω	$\triangleq$	Set of all the possible assignments of $X = \{X_i\}_{i\in[d_X]}$, $\Omega = \bigotimes_{i\in[d_X]}[N]$
n	$\triangleq$	Index on the support of random variables, $n \in [N]$
$\mathcal{H}_k$	$\triangleq$	History of observations $((x_{h,l}, r_{h,l}))_{h\in[H],l\in[k-1]}$ until episode k
Z_k	$\triangleq$	Random variable of the global factorization at episode k
Z_j^k	$\triangleq$	Random variable of the local factorization at episode k for factor j
Z_*	$\triangleq$	Random variable of the true factorization
Z_{j*}	$\triangleq$	Random variable of the true factorization for j-th factor
$\mathbb{E}_k[\cdot]$	$\triangleq$	Conditional expectation given history $\mathcal{H}_k$, $\mathbb{E}_k[\cdot] := \mathbb{E}[\cdot \mid \mathcal{H}_k]$
$\mathbb{P}_k[\cdot]$	$\triangleq$	Conditional probability given history $\mathcal{H}_k$, $\mathbb{P}_k[\cdot] := \mathbb{P}[\cdot \mid \mathcal{H}_k]$

B Parametric priors and posterior updates

In the following, we detail how the hyper-priors and priors of C-PSRL (Algorithm 1) can be specified through parametric distributions, and how the corresponding parameters are updated with the evidence provided by the collected data.

The hyper-prior $P_0 = \{P_{0,j}\}_{j=1}^{d_Y}$ is defined through d_Y distributions over the set of *local factorizations* $\mathcal{Z}_1, \ldots, \mathcal{Z}_{d_Y}$ resepctively, where each $\mathcal{Z}_j$ contains the parents assignments for the variable Y_j consistent with the graph prior $\mathcal{G}_0$. Let assume any arbitrary ordering of the local factorizations $z_i \in \mathcal{Z}_j$, such that each z_i is indexed by $i \in [|\mathcal{Z}_j|]$. Then, we can specify the hyper-prior j as a categorical distribution

$$P_{0,j}(z_i; \boldsymbol{\omega}) = \mathrm{Cat}(i; \boldsymbol{\omega}) = \frac{\omega_i}{\sum_t \omega_t},$$

where the sum is over $t \in [|\mathcal{Z}_j|]$, and the vector of parameters $\boldsymbol{\omega}$ is initialized as $\boldsymbol{\omega} = (1, \ldots, 1)$.

Then, for each local factorization $z_i \in \mathcal{Z}_j$ of the variable Y_j, we specify the prior $P_{0,j}(\cdot|z_i)$ over the model parameters of the corresponding transition factor p_j. The transition factor p_j is an $(N^{|\mathcal{Z}_j|}, N)$ stochastic matrix. The prior is specified through a Dirichlet distribution for each row of p_j, i.e.,

$$p_{0,j}(\cdot \mid z_i; \boldsymbol{\alpha}) = \mathrm{Dir}(\theta_1, \ldots, \theta_N; \alpha_1, \ldots, \alpha_N) = \frac{1}{B(\boldsymbol{\alpha})} \prod_{n=1}^{N} \theta_n^{\alpha_n - 1},$$

where $\boldsymbol{\alpha}$ is a vector of parameters initialized as $\boldsymbol{\alpha} = (1, \ldots, 1)$ and $B(\alpha) = \prod_{n=1}^{N} \Gamma(\alpha_n)/\Gamma(\sum_n \alpha_n)$ is a normalizing factor.

Having specified the hyper-prior and prior, we now show how to update them with the new evidence. Let be $\theta_1, \ldots, \theta_N \sim P_{k,j}(\cdot|x[z_j]; \boldsymbol{\alpha})$, and assume to collect the transition $(x[z_j], y[j] = i)$ from the true FMDP $\mathcal{F}_*$. Then, the posterior is

$$P_{k+1,j}(\theta_1, \ldots, \theta_N) \propto P(y[j] = i \mid \theta_1, \ldots, \theta_N) P_{k,j}(\theta_1, \ldots, \theta_N; \boldsymbol{\alpha}) \propto \theta_i \prod_{n=1}^{N} \theta^{\alpha_n - 1},$$

which is still a Dirichlet distribution with parameters $\mathrm{Dir}(\theta_1, \ldots, \theta_N; \alpha_1, \ldots, \alpha_i + 1, \ldots, \alpha_N)$. Then, we can propagate the posterior up the hierarchy to update the hyper-prior as

$$P_{k+1,j}(z) \propto P(y[j] = i \mid z) P_{t,j}(z; \boldsymbol{\omega})$$

$$\propto P_{k,j}(z; \boldsymbol{\omega}) \int P(y[j] = i \mid \theta_1, \ldots, \theta_N) P_{t,j}(\theta_1, \ldots, \theta_N \mid z) \mathrm{d}(\theta_1, \ldots, \theta_N) \tag{1}$$

$$\propto P_{k,j}(z; \boldsymbol{\omega}) \int \theta_i \frac{1}{B(\boldsymbol{\alpha})} \prod_{n=1}^{N} \theta_n^{\alpha_n - 1} \mathrm{d}(\theta_1, \ldots, \theta_N) \tag{2}$$

$$\propto P_{k,j}(z; \boldsymbol{\omega}) \frac{1}{B(\boldsymbol{\alpha})} B(\alpha_1, \ldots, \alpha_i + 1, \ldots, \alpha_N) \tag{3}$$

$$\propto P_{k,j}(z; \boldsymbol{\omega}) \frac{\alpha_i + 1}{\sum_t \alpha_t + 1} \tag{4}$$

where (2) is obtained by plugging the parametric prior in (1), we derive (3) by computing the integral over the simplex of $(\theta_1, \ldots, \theta_N)$, and (4) follows from $\Gamma(\alpha_i + 1) \prod_{t \neq i} \Gamma(\alpha_t) = (\alpha_i + 1) \prod_t \Gamma(\alpha_t)$ and $\Gamma(\alpha_i + 1 + \sum_{t \neq i} \alpha_t) = \Gamma(\sum_t \alpha_t) \sum_t (\alpha_t + 1)$. The resulting posterior is still a categorical distribution with the parameters $\omega_i \leftarrow \omega_i \frac{\alpha_i + 1}{\sum_t \alpha_t + 1}$.

C   Weak causal discovery

In the following, we show that we can extract, under a relatively mild causal minimality assumption, a Z-sparse super-DAG of the true causal graph $\mathcal{G}_{\mathcal{F}_*}$ as a byproduct of a run of Algorithm 1. We call this result *weak causal discovery*, to make a clear distinction between discovering a sparse super-DAG of a causal graph and true causal discovery, in which the minimal graph is discovered.

As required for any causal discovery algorithm, we need to state an assumption that connects the causal graph $\mathcal{G}_{\mathcal{F}_*}$ with the distribution p_* (i.e., the transition model) from which our observations are sampled in an i.i.d. manner [Spirtes et al., 2000]. Typically, in causal discovery, it is assumed that p_* fulfills the faithfulness assumption with regard to $\mathcal{G}_{\mathcal{F}_*}$, i.e., every independence in p_* implies a d-separation (see Definition 4 below) in $\mathcal{G}_{\mathcal{F}_*}$. Faithfulness, however, is a rather strong assumption which can be violated by path cancellations or xor-type dependencies, and weaker assumptions have been proposed [Spirtes et al., 2000, Pearl, 2009, Marx et al., 2021]. In this work, we build upon a strictly weaker assumption than faithfulness: causal minimality [Spirtes et al., 2000].[6]

Definition 3 (Causal Minimality). *A distribution P satisfies causal minimality with respect to a DAG $\mathcal{G}$ if P fulfills the Markov factorization property with respect to $\mathcal{G}$, but not with respect to any proper subgraph of $\mathcal{G}$.*

Intuition.   More intuitively, a distribution is minimal with respect to $\mathcal{G}$ if and only if there is no node that is conditionally independent of any of its parents, given the remaining parents [Peters et al., 2017]. There are two important points in this statement: i) none of the parents of a node is redundant, and ii) the dependence to a parent may only be detected given the remaining parents. Aspect ii) is a strictly weaker statement than required by faithfulness, which can be illustrated with a simple example. Consider the causal structure $X \to Y \leftarrow Z$, where all random variables are binary. If we generate X and Z via an unbiased coin and assign Y as $Y := X$ xor Z, Y will be marginally independent of X, as well as marginally independent of Z. However, Y is not independent of X (resp. Z) when we condition on its second parent Z (resp. X). Such an example violates faithfulness, i.e., there is a causal edge that is not matched by a dependence, but it does not violate causal minimality. For a more detailed discussion on such triples, we refer to [Marx et al., 2021].

In our context, Algorithm 1 has a positive probability of sampling all parents jointly (or a superset of them), and does not rely on checking pairs individually. Therefore, we can build upon the weaker assumption, causal minimality. Beyond identifiability in the limit, we are interested in the finite sample behaviour of our approach. Therefore, we propose a slightly stronger assumption for the value function, which is inspired by causal minimality.

Definition 2 (ϵ-Value Minimality). *An FMDP $\mathcal{F}$ fulfills ϵ-value minimality, if for any FMDP $\mathcal{F}'$ encoding a proper subgraph of $\mathcal{G}_{\mathcal{F}}$, i.e., $\mathcal{G}_{\mathcal{F}'} \subset \mathcal{G}_{\mathcal{F}}$, it holds that $V_{\mathcal{F}}^* > V_{\mathcal{F}'}^* + \epsilon$, where $V_{\mathcal{F}}^*$, $V_{\mathcal{F}'}^*$ are the value functions of the optimal policies in $\mathcal{F}$, $\mathcal{F}'$ respectively.*

Intuitively, ϵ-value minimality ensures that if we were to miss a true parent, the resulting optimal value function would be at most ϵ-optimal compared to the optimal value function evaluated on a graph that contains all true parents. Based on this rather lightweight assumption, we can extract from Algorithm 1 a graph $\mathcal{G}_{\mathcal{F}_K}$ that is guaranteed to be either the true DAG $\mathcal{G}_{\mathcal{F}_*}$, or a Z-sparse super-DAG of $\mathcal{G}_{\mathcal{F}_*}$ with high probability.

---

[6]The definition refers SGS-minimality proposed by Spirtes et al. [2000]. There exists an alternative definition called P-minimality, proposed by Pearl [2009]. In our setting, both assumptions are equivalent, since they only differ on graphs that violate triangle faithfulness [Zhang, 2020a, Zhang and Spirtes, 2008]. Since no nodes within $\mathcal{X}$ or within $\mathcal{Y}$ are allowed to be adjacent, such triangle structures cannot occur within our assumptions.

Corollary 5.1 (Weak Causal Discovery)**.** *Let $\mathcal{F}_*$ be an FMDP in which the transition model p_* fulfills the causal minimality assumption with respect to $\mathcal{G}_{\mathcal{F}_*}$, and let $\mathcal{F}_*$ fulfill ϵ-value minimality. Then, $\mathcal{G}_{\mathcal{F}_*} \subseteq \mathcal{G}_{\mathcal{F}_K}$ holds with high probability, where $\mathcal{G}_{\mathcal{F}_K}$ is a Z-sparse graph randomly selected within the sequence $\{\mathcal{G}_{\mathcal{F}_k}\}_{k=0}^{K-1}$ produced by C-PSRL over $K = \tilde{O}(H^5 d_Y^2 2^{d_X - \eta}/\epsilon^2)$ episodes.*

Proof. From Theorem 4.1, we have that the K-episodes Bayesian regret of Algorithm 1 is

$$\mathbb{E}\left[\sum_{k=0}^{K-1} V_*(\pi_*) - V_*(\pi_t)\right] \leq C_1 \sqrt{H^5 d_Y^2 2^{d_X - \eta}} \cdot \sqrt{K},$$

with high probability for some constant C_1 that does not depend on K. Through a standard regret-to-pac argument [Jin et al., 2018], it follows

$$\mathbb{E}\left[V_*(\pi_*) - V_*(\pi_K)\right] \leq C_2 \sqrt{H^5 d_Y^2 2^{d_X - \eta}} \cdot \frac{1}{\sqrt{K}} \tag{5}$$

with high probability for some constant C_2 that does not depend on K, and for a policy π_K that is randomly selected within the sequence of policies $\{\pi_k\}_{k=0}^{K-1}$ produced by Algorithm 1. By noting that π_K can be ϵ-optimal in the true FMDP $\mathcal{F}_*$ only if $\mathcal{G}_{\mathcal{F}_*} \subseteq \mathcal{G}_{\mathcal{F}_K}$ through the ϵ-value minimality assumption (Definition 2), we let $\mathbb{E}\left[V_*(\pi_*) - V_*(\pi_K)\right] = \epsilon$ in (5), which gives $K \geq C_2 H^5 d_Y^2 2^{d_X - \eta}/\epsilon^2$ and concludes the proof. $\qquad\square$

d-Separation. For the reader's convenience, here we report a brief definition of d-separation. More details can be found in [Peters et al., 2017].

Definition 4 (d-Separation)**.** *A path $\langle X, \ldots, Y \rangle$ between two vertices X, Y in a DAG is d-connecting given a set $\boldsymbol{Z}$, if*

1. *every collider[7] on the path is an ancestor of $\boldsymbol{Z}$, and*

2. *every non-collider on the path is not in $\boldsymbol{Z}$.*

If there is no path d-connecting X and Y given $\boldsymbol{Z}$, then X and Y are d-separated given $\boldsymbol{Z}$. Sets $\boldsymbol{X}$ and $\boldsymbol{Y}$ are d-separated given $\boldsymbol{Z}$, if for every pair X, Y, with $X \in \boldsymbol{X}$ and $Y \in \boldsymbol{Y}$, X and Y are d-separated given $\boldsymbol{Z}$.

---

[7]A collider C on a path $\langle \ldots, Q, C, W, \ldots \rangle$ is a node with two arrowhead pointing towards it, i.e. $\rightarrow C \leftarrow$.

D   Regret analysis

In this section, we provide the full derivation of the following result.

Theorem 4.1. *Let $\mathcal{G}_0$ be a causal graph prior with degree of sparseness Z and degree of prior knowledge η. The K-episodes Bayesian regret incurred by C-PSRL is*

$$\mathcal{BR}(K) = \tilde{O}\left(\left(H^{5/2}N^{1+Z/2}d_Y + \sqrt{H2^{d_X - \eta}}\right)\sqrt{K}\right).$$

On a high level, the proof is made up of two parts. The first part (presented in Section D.1) consists of decomposing the Bayesian regret into two components and then upper bounding the two expressions separately. This leads to the intermediate regret bound for a general latent hypothesis space, i.e., where the hypothesis space is not necessarily a product space, reported in Section D.1. The second part of the proof refines the analysis by considering a product latent hypothesis space (Section D.2) and the degree of prior knowledge (Section D.3), ultimately reaching the theorem statement.

We define the set $\Omega = \bigotimes_{i \in [d_X]}[N]$ of all the possible assignments of $X = \{X_i\}_{i \in [d_X]}$. For the sake of concision, we will denote $p_k\left(y[j] \mid x[z_j]\right)$ as $p_k\left(x[z_j]\right)$ where it will not lead to ambiguity. Moreover, we denote $\mathbb{E}_k[\cdot] := \mathbb{E}[\cdot \mid \mathcal{H}_k]$ and $\mathbb{P}_k[\cdot] := \mathbb{P}[\cdot \mid \mathcal{H}_k]$ the conditional expectation and probability given the history of observations $\mathcal{H}_k = ((x_{h,l}, r_{h,l}))_{h \in [H], l \in [k-1]}$ collected until episode k. Auxiliary results and lemmas mentioned alongside the analysis are reported in the Sections D.4 and D.5.

D.1   Analysis for a general latent hypothesis space

We first report a decomposition of the Bayesian regret and then proceeds to bound each component separately, which are then combined in a single regret rate.

Bayesian regret decomposition.   For episode k, we define $\overline{V}_k(\pi, z) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot \mid z)}[V_{\mathcal{F}}(\pi)]$ as the expected value of policy π according to the posterior conditioned on the latent factorization $z \sim P_k$ and history $\mathcal{H}_k$. As shown in [Russo and Van Roy, 2014, Proposition 1] for the bandit setting and in [Hong et al., 2022b, Section 5.1, Equation 6] for the reinforcement learning setting, we can decompose the Bayesian regret as

$$\mathcal{BR}(n) = \mathbb{E}\left[\sum_{k=1}^{K}\mathbb{E}_k\left[V_*(\pi_*) - \overline{V}_k(\pi_*, Z_*)\right]\right] + \mathbb{E}\left[\sum_{k=1}^{K}\mathbb{E}_k\left[\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\right]\right] \qquad (6)$$

by adding and subtracting $\overline{V}_k(\pi_*, Z_*)$ and noticing that π_*, Z_* are identically distributed to π_k, Z_k given $\mathcal{H}_k$. Notice that Z_k and Z_* indicate random variables, while we will indicate with the lowercase counterpart specific values of these random variables. The first term represents the regret incurred due to the concentration of the posteriors of the reward and transition models given the true factorization, while the second term captures the cost to identify the true latent factorization. We will bound each term of (6) separately.

Upper bounding the first term of (6).   For episode k, we define the event

$$E_k = \left\{\forall\, j \in [d_Y], \forall\, x[z_j] \in \Omega : \left|R_{\mathcal{F}_k}(x[z_j]) - \bar{r}_k(x[z_j^k])\right| \le c_k(x[z_j^k])\right.$$

$$\left. \text{and } \|p_{\mathcal{F}_k}(x[z_j]) - \bar{p}_k(x[z_j^k])\|_1 \le \phi(x[z_j^k]) \right\}$$

where the quantities are defined as follows. $\mathcal{F}_k$ denotes the FMDP sampled at episode k having mean reward $R_{\mathcal{F}_k}(x[z_j])$ and transition model $p_{\mathcal{F}_k}(x[z_j])$ for all $x[z_j] \in \Omega$. The expression

21

$\bar{r}_k(x[z_j]) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot|z)}[R_{\mathcal{F}}(x[z_j])]$ denotes the posterior mean of $R_{\mathcal{F}_k}(x[z_j])$, while $\bar{p}_k(x[z_j]) = (\bar{p}_k(y[j] = n \mid x[z_j]))_{n \in [N]}$ with $\bar{p}_k(y[j] = n \mid x[z_j]) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot|z)}[p_{\mathcal{F}}(y[j] = n \mid x[z_j])]$ denotes the posterior mean transition probability vector of size N for the j-th factor given a factorization z. With $c_k(x[z_j^k])$ and $\phi(x[z_j^k])$ we denote high-probability confidence widths for the j-th factor of the mean reward and transition model respectively. A detailed derivation of such confidence widths can be found in Section D.4. Informally, the event E_k expresses how close the mean rewards and transition models sampled at the episode k are to their posterior means. We refer with $\bar{E}_k$ to the complementary event of E_k.

Now, by reminding that π_*, Z_* are identically distributed to π_k, Z_k given $\mathcal{H}_k$, we can rewrite each element of the sum within the first term of (6) as

$$\mathbb{E}_k\left[V_{\mathcal{F}_k}(\pi_k) - \overline{V}_k(\pi_k, Z_k)\right]$$

$$\overset{(1)}{=} \mathbb{E}_k\left[\mathbb{E}_{\mathcal{F} \sim P_k(\cdot|Z_k)}\left[V_{\mathcal{F}_k}(\pi_k) - V_{\mathcal{F}}(\pi_k)\right]\right]$$

$$\overset{(2)}{\leq} \mathbb{E}_k\left[\sum_{h=1}^{H}\left(R_{\mathcal{F}_k}(S_{k,h}, A_{k,h}) - \bar{r}_k(S_{k,h}, A_{k,h}, Z_k)\right) + H\|p_{\mathcal{F}_k}(S_{k,h}, A_{k,h}) - \bar{p}_k(S_{k,h}, A_{k,h}, Z_k)\|_1\right]$$

$$\overset{(3)}{\leq} \mathbb{E}_k\left[\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\left(R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k])\right) + H\sum_{j=1}^{d_Y}\|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1\right]$$

$$\leq \mathbb{E}_k\left[H\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\left(R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k])\right) + \|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1\right]$$

$$\overset{(4)}{\leq} \mathbb{E}_k\left[H\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\left(R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k]) + \|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1\right)\mathbb{1}\{\bar{E}_k\}\right]$$

$$+ \mathbb{E}_k\left[H\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\left(c_k(X_{k,h}[Z_j^k]) + \phi_k(X_{k,h}[Z_j^k])\right)\mathbb{1}\{E_k\}\right] \tag{7}$$

where we have used the definition of $\overline{V}_k(\pi_k, Z_k)$ in step (1), Lemma D.3 in step (2), Lemma D.4 in step (3), and the definition of E_k in step (4).

By defining $\beta_k(X_{k,h}[Z_j^k]) := c_k(X_{k,h}[Z_j^k]) + \phi_k(X_{k,h}[Z_j^k])$ as the sum of both confidence widths, we can bound the second term of (7) by using Lemma D.5, while we bound the first term of the same equation by showing that $\bar{E}_k$ conditioned on $\mathcal{H}_k$ is unlikely. We rewrite the first term of (6) as

$$H\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\left(\mathbb{E}_k\left[\left(R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k])\right)\mathbb{1}\{\bar{E}_k\}\right]\right.$$

$$\left. + \mathbb{E}_k\left[\|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1\mathbb{1}\{\bar{E}_k\}\right]\right) \tag{8}$$

where we have distributed the indicator function. For the first term within the sums of (8), we have

$$\mathbb{E}_k\left[\left(R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k])\right)\mathbb{1}\{\bar{E}_k\}\right] \tag{9}$$

$$\leq \sum_{x[Z_j] \in \Omega}\int_{r=c_k(X_{k,h}[Z_j^k])}^{\infty} r\,\mathbb{P}_k\left((R_{\mathcal{F}_k}(x[Z_j]) - \bar{r}_k(x[Z_j^k])) = r\right)dr \tag{10}$$

$$\leq \sum_{x[Z_j] \in \Omega}\mathbb{P}_k\left(R_{\mathcal{F}_k}(x[Z_j]) - \bar{r}_k(x[Z_j^k]) \geq c_k(x[Z_j^k])\right) \tag{11}$$

$$\overset{(1)}{\leq} \sum_{x[Z_j]\in\Omega} \exp\left(-\frac{c_k(x[Z_j^k])^2}{2/4(\|\alpha_k^R(x[Z_j^k])\|_1+1)}\right) \tag{12}$$

$$\overset{(2)}{=} \sum_{x[Z_j]\in\Omega} \exp\left(-\frac{\frac{\log(2Kd_YN^Z)}{2(\|\alpha_k^R(x[Z_j^k])\|_1+1)}}{2/4(\|\alpha_k^R(x[Z_j^k])\|_1+1)}\right) \tag{13}$$

$$= \sum_{x[Z_j]\in\Omega} \exp\left(-\log(2Kd_YN^Z)\right) = \frac{1}{2Kd_Y} \tag{14}$$

In step (1) we have used Lemma D.2 and D.1, and in step (2) we have plugged-in the definition of $c_k(x[Z_j^k])$ from (plugged-in σ^2 of R_Δ), where $\alpha_k^R(x[Z_j^k])$ represents the parameters of the posterior over the mean reward for the j-th factor at episode k given factorization Z_j^k. For the second term within the sums of (8), we have

$$\underset{k}{\mathbb{E}}\left[\|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1 \mathbb{1}\{\bar{E}_k\}\right]$$

$$\overset{(1)}{\leq} N\underset{k}{\mathbb{E}}\left[\max_{n\in[N]}|p_{\mathcal{F}_k}(X_{k,h}[Z_j],n) - \bar{p}_k(X_{k,h}[Z_j^k],n)|\mathbb{1}\{\bar{E}_k\}\right]$$

$$\overset{(2)}{\leq} \sum_{x[Z_j]\in\Omega}\sum_{n\in[N]}\int_{p=\phi_k(X_{k,h}[Z_j^k])/\sqrt{N}}^{\infty} p\,\mathbb{P}_k(|p_{\mathcal{F}_k}(X_{k,h}[Z_j],n) - \bar{p}_k(X_{k,h}[Z_j^k],n)| = p)\,dp$$

$$\leq 2\mathbb{P}_k\left(|p_{\mathcal{F}_k}(X_{k,h}[Z_j],n) - \bar{p}_k(X_{k,h}[Z_j^k],n)| \geq \frac{\phi_k(x[Z_j^k])}{\sqrt{N}}\right)$$

$$\overset{(3)}{\leq} \sum_{x[Z_j]\in\Omega}\sum_{n\in[N]} 2\exp\left(-\frac{\phi_k(x[Z_j^k])^2}{2N/4(\|\alpha_k^P(x[Z_j^k])\|_1+1)}\right)$$

$$= \frac{N}{Kd_Y} \tag{15}$$

The steps are analogous to the ones for upper bounding the first term of (8). Specifically, in step (1) we use a trivial upper bound on the l^1-norm and in step (2) we divide the confidence width by the square root of the vector length $\sqrt{N}$ according to lemma [Lattimore and Szepesvári, 2020, Theorem 5.4.c]. The parameters $\alpha_k^P(x[Z_j^k])$ introduced in step (3) represent the parameters of the posterior over the transition model for the j-th factor at episode k given factorization Z_j^k.

By plugging (14) and (15) into (8) and then (8) into (7), we can bound the first term of (6) as

$$\mathbb{E}\left[\sum_{k=1}^{K}\underset{k}{\mathbb{E}}\left[V_*(\pi_*) - \overline{V}_k(\pi_*, Z_*)\right]\right] \leq NH^2 + H\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\underset{k}{\mathbb{E}}\left[\beta_k(X_{k,h}[Z_j^k])\right] \tag{16}$$

where we recall that $\beta_k(X_{k,h}[Z_j^k]) := c_k(X_{k,h}[Z_j^k]) + \phi_k(X_{k,h}[Z_j^k])$.

Upper bounding the second term of (6). Since there is no fundamental distinction between latent states in the tabular and factored MDP settings, our analysis in this section is aligned with [Hong et al., 2022b, Appendix B.3, step 2] and aims at effectively translating it into the factored MDPs notation.

In order to bound the second term of (6), we first need to define confidence sets over latent factorizations. For each episode k, we define a set of factorizations C_k so that $Z_* \in C_k$ with high probability. Since the latent factorization is unobserved, we can only exploit a proxy statistic for how well the model parameter posterior of each latent factorization predicts the rewards. We start defining a counting function $N_k(z) = \sum_{l=1}^{k-1}\mathbb{1}\{Z_l = z\}$ as the number of

times the factorization z has been sampled until episode k. Next, we define the following statistic associated with a factorization z and episode k,

$$G_k(z) = \sum_{l=1}^{k-1} \mathbb{1}\{Z_l = z\} \left(\overline{V}_l(\pi_l, z) - H\sqrt{2} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \beta_l(X_{l,h}[z_j]) - \sum_{h=0}^{H-1} \sum_{j=1}^{d_Y} R_{l,h}[j] \right)$$

The latter represents the total under-estimation of observed returns, as it expresses the difference between the lower confidence bound on the returns and the observed ones, assuming that z is the true latent factorization. Now we can define $C_k = \{z \in \mathcal{Z} : G_k(z) \leq \sqrt{HN_k(z)\log K}\}$ as the set of latent factorizations with at most $\sqrt{HN_k(z)\log K}$ excess. In the following, we show that $Z_* \in C_k$ holds with high probability for any episode.

Fix $Z_* = z$. Let $\mathcal{T}_{k,z} = \{l < k : Z_l = z\}$ the set of episodes where z has been sampled until episode k. We will first upper bound $G_k(z)$ by a martingale with respect to the history, then bound the martingale using Azuma-Hoeffding's inequality. We define the event

$$\mathcal{E}_{k,h,j} = \left\{ \left| R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k]) \right| \leq \sqrt{2}c_k(X_{k,h}[Z_j^k]) \right.$$

$$\left. \text{and } \|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1 \leq \sqrt{2}\phi_k(X_{k,h}[Z_j^k]) \right\}$$

in which the sampled reward and transition probabilities for factor j in step h of episode k are close to their posterior means. Let $\mathcal{E} = \cup_{k=1}^{K} \cup_{h=1}^{H} \cup_{j=1}^{d_Y} \mathcal{E}_{k,h,j}$ be the event that this holds for every factor, step, and episode. By union bound we have that

$$\mathbb{P}_k(\bar{\mathcal{E}}_{k,h,j}) \leq \mathbb{P}_k(\left| R_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{r}_k(X_{k,h}[Z_j^k]) \right| \geq \sqrt{2}c_k(X_{k,h}[Z_j^k]))$$

$$+ \mathbb{P}_k(\|p_{\mathcal{F}_k}(X_{k,h}[Z_j]) - \bar{p}_k(X_{k,h}[Z_j^k])\|_1 \geq \sqrt{2}\phi_k(X_{k,h}[Z_j^k]))$$

$$\overset{(1)}{\leq} \exp\left(\frac{2c_k(X_{k,h}[Z_j^k])^2}{\sigma^2} \right) + \exp\left(\frac{2\phi_k(X_{k,h}[Z_j^k])^2}{N\sigma^2} \right)$$

$$\leq (Kd_Y N^{d_X})^{-2}$$

where we have used Lemmas D.2 and D.1 in step (1) and $\bar{\mathcal{E}}_{k,h,j}$ is the complementary event of $\mathcal{E}_{k,h,j}$. Hence, for $\bar{\mathcal{E}} = \cup_{k=1}^{K} \cup_{h=1}^{H} \cup_{j=1}^{d_Y} \bar{\mathcal{E}}_{k,h,j}$, we have

$$\mathbb{P}(\bar{\mathcal{E}}) = \sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{z \in \mathcal{Z}} \sum_{j=1}^{d_Y} \sum_{x[Z_j] \in \Omega} \mathbb{P}_k(\bar{\mathcal{E}}_{k,h,j}) \leq \sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{z \in \mathcal{Z}} \sum_{j=1}^{d_Y} \sum_{x[Z_j] \in \Omega} (Kd_Y N^{d_X})^{-2} \leq H|\mathcal{Z}|K^{-1}$$

For episode $l \in \mathcal{T}_{k,z}$, let $\Delta_l = V_*(\pi_l) - \sum_{h=1}^{H} \sum_{j=1}^{d_Y} R_{l,h}[j]$. Since $\mathbb{E}_l[\Delta_l] = 0$, $(\Delta_l)_{l \in \mathcal{T}_{k,z}}$ is a martingale difference sequence with respect to the histories $(\mathcal{H}_l)_{l \in \mathcal{T}_{k,z}}$. Following exactly the same steps as in [Hong et al., 2022b, Proof of Lemma 7], we derive an upper bound on the probability of Z_* not being in the factorizations set C_k, namely:

$$\mathbb{P}(Z_* \notin C_k) \leq 2|\mathcal{Z}|HK^{-1} \tag{17}$$

We can now decompose the second term of (6) according to whether the sampled latent factorization is in C_k or not. Formally, we have

$$\mathbb{E}\left[\sum_{k=1}^{K} \mathbb{E}_k \left[\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k) \right] \right] \leq \mathbb{E}\left[\sum_{k=1}^{K} \left(\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k) \right) \mathbb{1}\{Z_k \in C_k\} \right]$$

$$+ H \sum_{k=1}^{K} \mathbb{P}(Z_* \notin C_k) \tag{18}$$

From the previous steps, using (17), we have that the second term of (18) is upper bounded by $2|\mathcal{Z}|H^2$, while in the following we derive an upper bound for the first term of (18) as in [Hong et al., 2022b, Appendix B.3, step 4]. We have

$$
\mathbb{E}\bigg[\sum_{k=1}^{K}\Big(\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\Big)\mathbb{1}\{Z_k \in C_k\}\bigg]
$$

$$
= H\sqrt{2}\,\mathbb{E}\bigg[\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\beta_k(X_{k,h}[Z_j^k])\bigg]
$$

$$
+ \mathbb{E}\bigg[\sum_{k=1}^{K}\Big(\overline{V}_k(\pi_k, Z_k) - H\sqrt{2}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\beta_k(X_{k,h}[Z_j^k]) - \sum_{h=1}^{H}\sum_{j=1}^{d_Y}R_{k,h}[j]\Big)\mathbb{1}\{Z_k \in C_k\}\bigg]
$$

$$
\overset{(1)}{\leq} H\sqrt{2}\,\mathbb{E}\bigg[\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\beta_k(X_{k,h}[Z_j^k])\bigg] + \mathbb{E}\bigg[\sum_{z\in\mathcal{Z}}G_{K+1}(z) + |\mathcal{Z}|H\bigg]
$$

$$
\overset{(2)}{\leq} H\sqrt{2}\,\mathbb{E}\bigg[\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\beta_k(X_{k,h}[Z_j^k])\bigg] + \sqrt{|\mathcal{Z}|KH\log K} + |\mathcal{Z}|H
$$

where in step (1) we use the definition of $G_{K+1}(z)$ and in step (2) we upper bound the same quantity.

Bayesian regret for a general latent hypothesis space. Combining the upper bounds of the two terms of (6), we get

$$
\mathcal{BR}(K) \leq NH^2 + H\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\mathbb{E}_k\big[\beta_k(X_{k,h}[Z_j^k])\big] + H\sqrt{2}\,\mathbb{E}\bigg[\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\beta_k(X_{k,h}[Z_j^k])\bigg]
$$

$$
+ \sqrt{|\mathcal{Z}|KH\log K} + |\mathcal{Z}|H + 2|\mathcal{Z}|H^2
$$

$$
\overset{(1)}{\leq} NH^2 + 3H^2d_YN^{d_X} + 3H^2d_YN\sqrt{N^{d_X}KH\log(4Kd_YN^{d_X})\log\Big(1 + \frac{KH}{2N^{d_X}\Lambda_{0,z}}\Big)}
$$

$$
+ \sqrt{|\mathcal{Z}|KH\log K} + 3|\mathcal{Z}|H^2
$$

where in step (1) we have used Lemma D.5, and we denote

$$
\Lambda_{0,z} = \min\bigg\{\min_{j,x[z_j]}\|\alpha_0^R(x[z_j])\|_1, \min_{j,x[z_j]}\|\alpha_0^P(x[z_j])\|_1\bigg\}
$$

Due to the Z-sparseness assumption, we can rewrite the Bayesian regret as

$$
\mathcal{BR}(K) \leq NH^2 + 3H^2d_YN^Z + 3H^2d_YN\sqrt{N^ZKH\log(4Kd_YN^Z)log\Big(1 + \frac{KH}{2N^Z\Lambda_{0,z}}\Big)}
$$

$$
+ \sqrt{|\mathcal{Z}|KH\log K} + 3|\mathcal{Z}|H^2
$$

$$
= \tilde{O}\Big(H^2d_YN^Z + H^{\frac{5}{2}}d_YN^{1+\frac{Z}{2}}\sqrt{K} + \sqrt{|\mathcal{Z}|KH} + |\mathcal{Z}|H^2\Big)
$$

Notably, this rate is sublinear in the number of episodes K and latent factorization $|\mathcal{Z}|$, exponential in the degree of sparseness Z.

D.2 Refinement 1: Product latent hypothesis spaces

As we briefly explained in Section 3, C-PSRL samples the factorization z from the product space $\mathcal{Z} = \mathcal{Z}_1 \times \ldots \times \mathcal{Z}_{d_Y}$ by combining independent samples $z_j \in \mathcal{Z}_j$ for each variable

Y_j. This allows us to refine the dependence in $|\mathcal{Z}|$ to $\overline{C} := \max_{j \in [d_Y]} |\mathcal{Z}_j| \leq |\mathcal{Z}|$. We can replicate the same steps of the previous section in order to derive the Bayesian regret for the setting with a product latent hypothesis space. For the sake of clarity, we report here the main steps highlighting the difference with the previous section.

For an episode $k \in [K]$, we define $C_k^j = \left\{ z_j \in \mathcal{Z}_j : G_k^j(\bar{z}) \leq \sqrt{H N_k^j(z_j) \log K} \right\}$ where

$$
G_k^j(z_j) = \sum_{l=1}^{k-1} \mathbb{1}\{Z_j^l = z_j\} H \sum_{h=1}^{H} \Big(\big(\bar{r}_l(X_{l,h}[z_j]) - R_{l,h}[j] \big)
$$
$$
+ \|\bar{p}_k(X_{l,h}[z_j]) - p_*(X_{l,h}[z_{j*}])\|_1 - \sqrt{2}\beta_k(X_{l,h}[z_j]) \Big)
$$

and $N_k^j(z_j) = \sum_{l=1}^{k-1} \mathbb{1}\{Z_j^l = z_j\}$. While $G_k^j(z_j)$ captures the under-estimation of the observed returns at the level of a single factor, $N_k^j(z_j)$ counts the number of times that the local factorization z_j has been sampled for node j until episode k. Next, we define $\mathcal{T}_{k,z_j}^j = \{l < k : Z_j^l = z_j\}$ as the set of episodes where z_j has been sampled for node j.

First, we can derive an upper bound of $\mathbb{P}(\bar{\mathcal{E}})$ depending on $\overline{C}$ by noticing that the inner-most sum depends only on the local factorization hence we can swap the two preceding sums over $\mathcal{Z}$ and d_Y as shown in step (1) of the following:

$$
\mathbb{P}(\bar{\mathcal{E}}) = \sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{z \in \mathcal{Z}} \sum_{j=1}^{d_Y} \sum_{x[z_j] \in \Omega} \mathbb{P}_k(\bar{\mathcal{E}}_{k,h,j})
$$
$$
\leq \sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{z \in \mathcal{Z}} \sum_{j=1}^{d_Y} \sum_{x[z_j] \in \Omega} (K d_Y N^{d_X})^{-2}
$$
$$
\overset{(1)}{=} \leq \sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \sum_{z_j \in \mathcal{Z}_j} \sum_{x[z_j] \in \Omega} (K d_Y N^{d_X})^{-2}
$$
$$
\leq H \overline{C} K^{-1}
$$

Now, we wish to upper bound $\mathbb{P}(Z_{j*} \notin C_k^j)$. For episode $l \in \mathcal{T}_{k,z_j}^j$ let $\Delta_l^j = \sum_{h=1}^{H} R_*(X_{l,h}[Z_{j*}]) - \sum_{h=1}^{H} R_{l,h}[j]$. Since $\mathbb{E}_l[\Delta_l^j] = 0$ we have that $(\Delta_l^j)_{l \in \mathcal{T}_{t,z_j}^j}$ is a martingale difference sequence with respect to the histories $(\mathcal{H}_l)_{l \in \mathcal{T}_{t,z_j}^j}$.

By following the same steps as in Hong et al. [2022b], we get

$$
G_k^j(z_j)\mathbb{1}\{\mathcal{E}\} = \sum_{l \in \mathcal{T}_{t,z_j}^j} H \sum_{h=1}^{H} \Big(\big(\bar{r}_l(X_{l,h}[z_j]) - R_{l,h}[j] \big)
$$
$$
+ \|\bar{p}_k(X_{l,h}[z_j]) - p_*(X_{l,h}[z_{j*}])\|_1 - \sqrt{2}\beta_k(X_{l,h}[z_j]) \Big) \leq \sum_{l \in \mathcal{T}_{k,z_j}^j} \Delta_l^j
$$

and by fixing $|\mathcal{T}_{t,z_j}^j| = N_t^j(z_j) = u < t$, and using Azuma-Hoeffding's inequality, we derive

$$
\mathbb{P}_k \left(G_k^j(\bar{z})\mathbb{1}\{\mathcal{E}\} \geq \sqrt{Hu \log K} \right) \leq \mathbb{P} \left(\sum_{l \in \mathcal{T}_{k,z_j}^j} \Delta_l^j \geq \sqrt{Hu \log K} \right) \leq K^{-2}.
$$

Therefore, by using union bounds, we can write

$$
\mathbb{P}(Z_* \notin C_k) \leq \sum_{j=1}^{d_Y} \mathbb{P}(Z_{j*} \notin C_k^j)
$$

$$\leq \sum_{j=1}^{d_Y} \sum_{z_j \in \mathcal{Z}_j} \sum_{u=1}^{k-1} \mathbb{P}\left(G_k^j(z_j) \geq \sqrt{Hu \log K}\right)$$

$$\leq \mathbb{P}(\bar{\mathcal{E}}) + \sum_{j=1}^{d_Y} \sum_{z_j \in \mathcal{Z}_j} \sum_{u=1}^{k-1} \mathbb{P}\left(G_k^j(z_j)\mathbb{1}\{\mathcal{E}\} \geq \sqrt{Hu \log K}\right)$$

$$\leq d_Y H \overline{C} K^{-1} \tag{19}$$

We can now decompose the second term of (6) according to whether the sampled latent factorization is in C_k or not, as in the previous section. Formally, we have

$$\mathbb{E}\left[\sum_{k=1}^{K} \mathbb{E}\left[\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\right]\right] \leq \mathbb{E}\left[\sum_{k=1}^{K} \left(\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\right)\mathbb{1}\{Z_k \in C_k\}\right]$$

$$+ H \sum_{k=1}^{K} \mathbb{P}(Z_* \notin C_k)$$

From (19), we know that the second term is upper bounded by $d_Y \overline{C} H^2$. Meanwhile, we can bound the first term as follows

$$\mathbb{E}\left[\sum_{k=1}^{K} \left(\overline{V}_k(\pi_k, Z_k) - V_*(\pi_k)\right)\mathbb{1}\{Z_k \in C_k\}\right]$$

$$\leq \mathbb{E}\left[\sum_{k=1}^{K} H \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \left(\bar{r}_k(X_{k,h}[Z_j^k]) - R_*(X_{k,h}[Z_{j*}])\right.\right.$$

$$\left.\left. + \|\bar{p}_k(X_{l,h}[Z_j^k]) - p_*(X_{l,h}[Z_{j*}])\|_1\right)\mathbb{1}\{Z_j^k \in C_k^j\}\right]$$

$$= H\sqrt{2}\,\mathbb{E}\left[\sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \beta_k(X_{k,h}[Z_j^k])\right] + \mathbb{E}\left[\sum_{k=1}^{K} H \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \left(\bar{r}_k(X_{k,h}[Z_j^k]) - R_*(X_{k,h}[Z_{j*}])\right.\right.$$

$$\left.\left. + \|\bar{p}_k(X_{l,h}[Z_j^k]) - p_*(X_{l,h}[Z_{j*}])\|_1 - \sqrt{2}\beta_k(X_{k,h}[Z_j^k])\right)\mathbb{1}\{Z_j^k \in C_k^j\}\right]$$

$$\overset{(1)}{\leq} H\sqrt{2}\,\mathbb{E}\left[\sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \beta_k(X_{k,h}[Z_j^k])\right] + \mathbb{E}\left[\sum_{j=1}^{d_Y} \sum_{z_j \in \mathcal{Z}_j} G_{K+1}^j(z_j) + d_Y \overline{C} H\right]$$

$$\overset{(2)}{\leq} H\sqrt{2}\,\mathbb{E}\left[\sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \beta_k(X_{k,h}[Z_j^k])\right] + \sum_{j=1}^{d_Y} \sum_{z_j \in \mathcal{Z}_j} \frac{1}{d_Y}\sqrt{HN_{K+1}^j(z_j) \log K} + d_Y \overline{C} H$$

$$\leq H\sqrt{2}\,\mathbb{E}\left[\sum_{k=1}^{K} \sum_{h=1}^{H} \sum_{j=1}^{d_Y} \beta_k(X_{k,h}[Z_j^k])\right] + \sqrt{\overline{C}KH \log K} + d_Y \overline{C} H$$

where in step (1) we use the definition of $G_{K+1}^j(z)$ and in step (2) we upper bound the same quantity.

Bayesian regret for a product latent hypothesis space. Exploiting the Z-sparseness assumption, we can write

$$\mathcal{BR}(K) = \tilde{O}\left(H^2 d_Y N^Z + H^{\frac{5}{2}} d_Y N^{1+\frac{Z}{2}} \sqrt{K} + \sqrt{\overline{C}KH} + d_Y \overline{C} H^2\right) \tag{20}$$

Notably, this rate is sublinear in the number of episodes K and the number of latent local factorizations $\overline{C}$, exponential in the degree of sparseness Z.

D.3 Refinement 2: Degree of prior knowledge

Finally, we aim to capture the dependency in the degree of prior knowledge η in the Bayesian regret. To do that, we have to express $\overline{C} = \max_{j \in [d_Y]} |\mathcal{Z}_j|$ in terms of η. We can write

$$\overline{C} = \sum_{i=0}^{Z} \binom{d_X}{i} = \sum_{i=0}^{Z} C_i^{d_X}$$

where we count the empty factorization when $i = 0$. Given a graph hyper-prior that fixes $\eta < Z$ edges for each node $j \in [d_Y]$, we can count the number of admissible local factorizations as

$$\overline{C} = \sum_{i=0}^{Z-\eta} \binom{d_X - \eta}{i}$$

where we count the factorization with only the edges fixed a priori when $i = 0$. We can build an upper bound on $\overline{C}$ as follows.

$$\begin{aligned}
\overline{C} &= \sum_{i=0}^{Z-\eta} \binom{d_X - \eta}{i} \\
&\leq 2^{d_X - \eta - 1} \exp\left(\frac{(d_X + \eta - 2Z - 2)^2}{4(1 + Z - d_X)}\right) \\
&\leq 2^{d_X - \eta} \exp\left(\frac{(d_X + \eta - 2Z)^2}{4(1 + Z - d_X)}\right) =: \phi(d_X, Z, \eta)
\end{aligned}$$

Since it is hard to interpret the rate of the latter upper bound, we derive a looser version that is easier to interpret. We have

$$\begin{aligned}
\overline{C} &= \sum_{i=0}^{Z-\eta} \binom{d_X - \eta}{i} \\
&= 2^{d_X - \eta} - \sum_{i=0}^{d_X - Z - 1} \binom{d_X - \eta}{i} \\
&\leq 2^{d_X - \eta} - 2^{d_X - Z} + 1 \\
&\leq 2^{d_X - \eta}
\end{aligned}$$

From the latter we can notice that each unit of the degree of prior knowledge η make the hypothesis space shrink with an exponential rate, and thus the corresponding regret terms as well. In particular, by plugging-in the upper bound $\overline{C} \leq 2^{d_X - \eta} = \frac{2^{d_X}}{2^\eta}$ in the Bayesian regret in (20), we obtain the final upper bound, which is

$$\mathcal{BR}(K) = \tilde{O}\left(H^2 d_Y N^Z + H^{\frac{5}{2}} d_Y N^{1 + \frac{Z}{2}} \sqrt{K} + \sqrt{2^{d_X - \eta} K H} + d_Y 2^{d_X - \eta} H^2\right). \quad (21)$$

D.4 High probability confidence widths

Here we define high-probability confidence widths on the reward function and transition model along the lines of Hong et al. [2022b], but with the difference that the confidence widths are defined for all factors and their possible assignments rather than for state-action pairs as in the tabular setting. We denote as $c_k(x[z_j])$ and $\phi_k(x[z_j])$ the confidence widths for the j-th factor of the reward function and the transition model respectively. In the following, we indicate with $R_{\mathcal{F}}$ and $p_{\mathcal{F}}$ the mean reward and transition model of the FMDP $\mathcal{F}$ respectively.

Reward function. First, we write the posterior mean reward for the j-th factor, given a factorization z as $\bar{r}_k(x[z_j]) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot|z)}[R_{\mathcal{F}}(x[z_j])]$. We wish to have a high probability bound of the type

$$\mathbb{P}_k \left(|R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j])| \geq c_k(x[z_j]) \right) \leq \frac{1}{K}$$

for all $j \in [d_Y]$ and possible assignments $x[z_j] \in \Omega$. By the union bound, we have

$$\mathbb{P}_k \left(\left| R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j]) \right| \geq c_k(x[z_j]) \right)$$
$$= \mathbb{P}_k \left(\bigcup_{j=1}^{d_Y} \bigcup_{x[z_j] \in \Omega} \left\{ |R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j])| \geq c_k(x[z_j]) \right\} \right)$$
$$\leq \sum_{j=1}^{d_Y} \sum_{x[z_j] \in \Omega} \mathbb{P}_k \left(|R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j])| \geq c_k(x[z_j]) \right).$$

Applying a union bound again to the latter expression, we can derive the following one-sided bound:

$$\mathbb{P}_k \left(R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j]) \geq c_k(x[z_j]) \right) \leq \frac{1}{2K d_Y N^{d_X}}.$$

According to Lemma D.1, $R_\Delta := R_{\mathcal{F}}(x[z_j]) - \bar{r}_k(x[z_j])$ is a σ^2-subgaussian random variable with $\sigma^2 = 1/\left(4 \left(\|\alpha_k^R(z_j)\|_1 + 1 \right) \right)$. Therefore, through the Cramèr-Chernoff method exploited in Lemma D.2, we have that the high probability bound above holds if

$$\exp \left(-\frac{c_k(x[z_j])^2}{2\sigma^2} \right) \leq \frac{1}{2K d_Y N^{d_X}}$$

which holds if and only if

$$c_k(x[z_j]) \geq \sqrt{2\sigma^2 \log(2K d_Y N^{d_X})}$$
$$= \sqrt{\frac{\log(2K d_Y N^{d_X})}{2 \left(\|\alpha_k^R(x[z_j])\|_1 + 1 \right)}} \qquad \text{(plugged-in } \sigma^2 \text{ of } R_\Delta\text{)}$$

Hence, we pick $c_k(x[z_j]) := \sqrt{\frac{\log(2K d_Y N^Z)}{2\left(\|\alpha_k^R(x[z_j])\|_1 + 1 \right)}}$, where Z is a lower bound on the value of d_X, which holds due to the Z-sparseness assumption.

Transition model. The derivation is analogous to the one for the reward function, hence here we report only the differences. First, we write the posterior mean transition probability for the j-th factor, given a factorization z as $\bar{p}_k(y[j] = n \mid x[z_j]) = \mathbb{E}_{\mathcal{F} \sim P_k(\cdot|z)}[p_{\mathcal{F}}(y[j] = n \mid x[z_j])]$, which is the probability, according to the posterior at time k, over the element $n \in [N]$ of the domain of the j-th component of Y. Since we want to bound the deviations over all components of a factor, we define the vector form of the previous expression as $\bar{p}_k(x[z_j]) = (\bar{p}_k(y[j] = n \mid x[z_j]))_{n \in [N]}$. By following the same steps as for the confidence width of the reward function, we get

$$\phi_k(x[z_j]) := \sqrt{\frac{N \log(2K d_Y N^Z)}{2 \left(\|\alpha_k^P(x[z_j])\|_1 + 1 \right)}} \qquad (22)$$

where the $\sqrt{N}$ term is due to [Lattimore and Szepesvári, 2020, Theorem 5.4.c], since the l^1 norm sums over $N\sigma^2$-subgaussian random variables and therefore the induced random variable is $(N\sigma^2)$-subgaussian.

D.5 Auxiliary lemmas

Lemma D.1 (Theorem 1 and 3 of Marchal and Arbel [2017]). *Let $X \sim \text{Beta}(\alpha, \beta)$ for $\alpha, \beta > 0$. Then $X - \mathbb{E}[X]$ is σ^2-subgaussian with $\sigma^2 = 1/(4(\alpha + \beta + 1))$. Similarly, let $X \sim \text{Dir}(\alpha)$ for $\alpha \in \mathbb{R}_+^d$. Then $X - \mathbb{E}[X]$ is σ^2-subgaussian with $\sigma^2 = 1/(4(\|\alpha\|_1 + 1))$.*

Lemma D.2 (Theorem 5.3 of Lattimore and Szepesvári [2020]). *If X is σ^2-subgaussian, then for any $\varepsilon \geq 0$,*

$$\mathbb{P}(X \geq \varepsilon) \leq \exp\left(-\frac{\varepsilon^2}{2\sigma^2}\right)$$

Lemma D.3 (Value Difference Lemma, Lemma 6 Hong et al. [2022b]). *For any MDPs $\mathcal{M}', \mathcal{M}$, and policy π,*

$$V_{\mathcal{M}'}(\pi) - V_{\mathcal{M}}(\pi) \leq \mathbb{E}_{\mathcal{M}}\left[\sum_{h=1}^{H} R_{\mathcal{M}'}(S_h, A_h) - R_{\mathcal{M}}(S_h, A_h) + H\|p_{\mathcal{M}'}(S_h, A_h) - p_{\mathcal{M}}(S_h, A_h)\|_1\right]$$

Proof. This upper bound can be obtained by trivially upper bounding with 1 the reward at each step, and therefore with h the value function within the statement in [Jin et al., 2020a, Lemma C.1]. □

Lemma D.4 (Deviations of Factored Reward and Transitions Osband and Van Roy [2014]). *Given two reward functions R and $\bar{R}$ with scopes $\{z_j\}_{j=1}^{d_Y}$ we can upper bound the deviations by*

$$|R(x) - \bar{R}(x)| \leq \sum_{j=1}^{d_Y} |R_j(x[z_j]) - \bar{R}_j(x[z_j])|$$

and, given two transition models p and $\bar{p}$ with scopes $\{z_j\}_{j=1}^{d_Y}$ we can upper bound the deviations by

$$|p(x) - \bar{p}(x)| \leq \sum_{j=1}^{d_Y} |p_j(x[z_j]) - \bar{p}_j(x[z_j])|$$

Lemma D.5. *For episode k and latent factorization $z \in \mathcal{Z}$, let $\beta_k(x[z_j]) = c_k(x[z_j]) + \phi_k(x[z_j])$ for any $j \in [d_Y]$ and $x[z_j] \in \Omega$. Let $\Lambda_{0,z} = \min\left\{\min_{j,x[z_j]}\|\alpha_0^R(x[z_j])\|_1, \min_{j,x[z_j]}\|\alpha_0^P(x[z_j])\|_1\right\}$ indicates the minimum level of concentration between the reward function and transition model priors for any factor and latent factorization z. Then, we have*

$$\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \beta_k(X_{k,h}[z_j]) \leq Hd_Y N^{d_X} + d_Y NH\sqrt{N^{d_X}KH\log(4Kd_Y N^{d_X})\log\left(1 + \frac{KH}{2N^{d_X}\Lambda_{0,z}}\right)}$$

Proof. We define $N_k(x[z_j]) = \sum_{l=1}^{k-1}\sum_{h=1}^{H} \mathbb{1}\{X_{l,h}[z_j] = x[z_j]\}$ as the number of times the assignment $x[z_j]$ was sampled up to episode k for factor j. We can decompose the sum as

$$\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \beta_k(X_{k,h}[z_j])$$

$$\leq \sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \mathbb{1}\{N_k(X_{k,h}[z_j]) \leq h\} + \sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \mathbb{1}\{N_k(X_{k,h}[Z_j]) > H\}\beta_k(X_{k,h}[z_j])$$

where we upper bound by 1 the regret in a step due to one factor. Therefore, the first term is upper bounded by $Hd_Y N^{d_X}$ since there are at most N^{d_X} assignments for each $j \in [d_Y]$ and the same one can appear in the sum at most H times, thus removing the dependency on K. Due to the assumption of Z-sparseness, we will later use the bound $Hd_Y N^{\bar{Z}}$. As for the

second term, we define $N_{k,h}(x[z_j]) = N_k(x[z_j]) + \sum_{p=1}^{h-1} \mathbb{1}\{X_{k,p}[z_j] = x[z_j]\}$ as the number of times $x[z_j]$ was sampled up to step h of episode k, for factor j. We split β_k into c_k and ϕ_k. For c_k we have:

$$
\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \mathbb{1}\{N_k(X_{k,h}[z_j]) > H\}c_k(X_{k,h}[z_j])
$$

$$
= \sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \mathbb{1}\{N_k(X_{k,h}[z_j]) > H\}\sqrt{\frac{\log(2Kd_Y N^Z)}{2\left(\|\alpha_k^R(X_{k,h}[z_j])\|_1 + 1\right)}}
$$

$$
= \sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\sum_{x[z_j]\in\Omega} \mathbb{1}\{N_k(x[z_j]) > H\}\sqrt{\frac{\log(2Kd_Y N^Z)}{2\|2\alpha_k^R(X_{k,h}[z_j])\|_1 + 2N_k(x[z_j]) + 2}}
$$

$$
\leq \sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y}\sum_{x[z_j]\in\Omega} \sqrt{\frac{\log(2Kd_Y N^Z)}{2\|2\alpha_k^R(X_{k,h}[z_j])\|_1 + N_{k,h}(x[z_j])}}
$$

$$
\leq \sqrt{\log(2Kd_Y N^Z)}\sum_{j=1}^{d_Y}\sum_{x[z_j]\in\Omega} \sqrt{N_{K+1}(x[z_j])\sum_{u=1}^{N_{K+1}(x[z_j])}\frac{1}{2\|\alpha_k^R(X_{k,h}[z_j])\|_1 + u}}
$$

$$
\leq \sqrt{N^{d_X}KH\log(2Kd_Y N^Z)}\sum_{j=1}^{d_Y}\sqrt{\sum_{u=1}^{KH/N^{d_X}}\frac{1}{2\Lambda_{0,z} + u}}
$$

$$
\leq d_Y\sqrt{N^{d_X}KH\log\left(2Kd_Y N^Z\right)\log\left(1 + \frac{KH}{2N^{d_X}\Lambda_{0,z}}\right)}
$$

where in the first step we have plugged-in c_k as picked in Section 4, in the second step have exploited the posterior update rule, in step three we have used that if $N_k(x[z_j]) > H$ we have that $N_{k,h}(x[z_j]) \leq N_k(x[z_j]) + H \leq 2N_k(x[z_j])$, and in the remaining passages we have used known bounds as in [Hong et al., 2022b, Lemma 6]. Analogously, for ϕ_k we can derive the following.

$$
\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \mathbb{1}\{N_k(X_{k,h}[z_j]) > H\}\phi_k(X_{k,h}[z_j])
$$

$$
\leq d_Y NH\sqrt{N^{d_X}KH\log\left(4Kd_Y N^Z\right)\log\left(1 + \frac{KH}{2N^{d_X}\Lambda_{0,z}}\right)}
$$

Notice that again, due to the Z-sparseness, we can replace in the steps above N^{d_X} with N^Z. Combining the terms, we have:

$$
\sum_{k=1}^{K}\sum_{h=1}^{H}\sum_{j=1}^{d_Y} \beta_k(X_{k,h}[z_j]) \leq Hd_Y N^{d_X} + 2d_Y NH\sqrt{N^{d_X}KH\log\left(4Kd_Y N^Z\right)\log\left(1 + \frac{KH}{2N^{d_X}\Lambda_{0,z}}\right)}
$$

$\square$